Package 'vcd'

August 3, 2014

Version 1.3-2

Date 2014-8-3

Title Visualizing Categorical Data

Description Visualization techniques, data sets, summary and inference
procedures aimed particularly at categorical data. Special
emphasis is given to highly extensible grid graphics. The
package was inspired by the book ``Visualizing Categorical Data'' by Michael Friendly.

LazyLoad yes

LazyData yes

Depends R (>= 2.4.0), grid, stats

Suggests KernSmooth, mvtnorm, kernlab, HSAUR, coin

Imports utils, MASS, grDevices, colorspace

License GPL-2

Author David Meyer [aut, cre],Achim Zeileis [aut],Kurt Hornik [aut],Florian Gerber [ctb],Michael Friendly [ctb]

Maintainer David Meyer <David.Meyer@R-project.org>

NeedsCompilation no

Repository CRAN

Date/Publication 2014-08-03 17:35:14

R **topics documented:**

---

agreementplot *Bangdiwala's Observer Agreement Chart*

Description

Representation of a $k \times k$ confusion matrix, where the observed and expected diagonal elements are represented by superposed black and white rectangles, respectively. The function also computes a statistic measuring the strength of agreement (relation of respective area sums).

Usage

```
## Default S3 method:
agreementplot(x, reverse_y = TRUE, main = NULL,
              weights = c(1, 1 - 1/(ncol(x) - 1)^2), margins = par("mar"),
              newpage = TRUE, pop = TRUE, xlab = names(dimnames(x))[2],
              ylab = names(dimnames(x))[1],
```

```
               xlab_rot = 0, xlab_just = "center",
               ylab_rot = 90, ylab_just = "center",
               fill_col = function(j) gray((1 - (weights[j]) ^ 2) ^ 0.5),
               line_col = "red", xscale = TRUE, yscale = TRUE,
                ...)
## S3 method for class 'formula'
agreementplot(formula, data = NULL, ..., subset)
```

Arguments

x	a confusion matrix, i.e., a table with equal-sized dimensions.
reverse_y	if TRUE, the y axis is reversed (i.e., the rectangles' positions correspond to the contingency table).
main	user-specified main title.
weights	vector of weights for successive larger observed areas, used in the agreement strength statistic, and also for the shading. The first element should be 1.
margins	vector of margins (see [par](#)).
newpage	logical; if TRUE, the plot is drawn on a new page.
pop	logical; if TRUE, all newly generated viewports are popped after plotting.
xlab, ylab	labels of x- and y-axis.
xlab_rot, ylab_rot	
	rotation angle for the category labels.
xlab_just, ylab_just	
	justification for the category labels.
fill_col	a function, giving the fill colors used for exact and partial agreement
line_col	color used for the diagonal reference line
formula	a formula, such as y ~ x. For details, see [xtabs](#).
data	a data frame (or list), or a contingency table from which the variables in formula should be taken.
subset	an optional vector specifying a subset of the rows in the data frame to be used for plotting.
xscale, yscale	logicals indicating whether the marginals should be added on the x-axis/y-axis, respectively.
...	further graphics parameters (see [par](#)).

Details

Weights can be specified to allow for partial agreement, taking into account contributions from off-diagonal cells. Partial agreement is typically represented in the display by lighter shading, as given by fill_col(j), corresponding to weights[j].

A weight vector of length 1 means strict agreement only, each additional element increases the maximum number of disagreement steps.

Value

Invisibly returned, a list with components

```
Bangdiwala          the unweighted agreement strength statistic.
Bangdiwala_Weighted
                    the weighted statistic.
weights             the weight vector used.
```

Author(s)

David Meyer <David.Meyer@R-project.org>

References

Bangdiwala, S. I. (1988). The Agreement Chart. Department of Biostatistics, University of North Carolina at Chapel Hill, Institute of Statistics Mimeo Series No. 1859, http://www.stat.ncsu.edu/information/library/mimeo.archive/ISMS_1988_1859.pdf

Bangdiwala, S. I., Ana S. Haedo, Marcela L. Natal, and Andres Villaveces. The agreement chart as an alternative to the receiver-operating characteristic curve for diagnostic tests. *Journal of Clinical Epidemiology*, 61 (9), 866-874.

Michael Friendly (2000), *Visualizing Categorical Data*. SAS Institute, Cary, NC.

Examples

```
data("SexualFun")
agreementplot(t(SexualFun))

data("MSPatients")
## best visualized using a resized device, e.g. using:
## get(getOption("device"))(width = 12)
pushViewport(viewport(layout = grid.layout(ncol = 2)))
pushViewport(viewport(layout.pos.col = 1))
agreementplot(t(MSPatients[,,1]), main = "Winnipeg Patients",
              newpage = FALSE)
popViewport()
pushViewport(viewport(layout.pos.col = 2))
agreementplot(t(MSPatients[,,2]), main = "New Orleans Patients",
              newpage = FALSE)
popViewport(2)
dev.off()
```

Arthritis   *Arthritis Treatment Data*

Description

Data from Koch \& Edwards (1988) from a double-blind clinical trial investigating a new treatment for rheumatoid arthritis.

Usage

```
data("Arthritis")
```

Format

A data frame with 84 observations and 5 variables.

ID patient ID.

Treatment factor indicating treatment (Placebo, Treated).

Sex factor indicating sex (Female, Male).

Age age of patient.

Improved ordered factor indicating treatment outcome (None, Some, Marked).

Source

Michael Friendly (2000), Visualizing Categorical Data: http://euclid.psych.yorku.ca/ftp/sas/vcd/catdata/arthrit.sas

References

G. Koch \& S. Edwards (1988), Clinical efficiency trials with categorical data. In K. E. Peace (ed.), *Biopharmaceutical Statistics for Drug Development*, 403–451. Marcel Dekker, New York.

M. Friendly (2000), *Visualizing Categorical Data*. SAS Institute, Cary, NC.

Examples

```
data("Arthritis")
art <- xtabs(~ Treatment + Improved, data = Arthritis, subset = Sex == "Female")
art

mosaic(art, gp = shading_Friendly)
mosaic(art, gp = shading_max)
```

---

assoc   *Extended Association Plots*

Description

Produce an association plot indicating deviations from a specified independence model in a possibly high-dimensional contingency table.

Usage

```
## Default S3 method:
assoc(x, row_vars = NULL, col_vars = NULL, compress = TRUE,
  xlim = NULL, ylim = NULL,
  spacing = spacing_conditional(sp = 0), spacing_args = list(),
  split_vertical = NULL, keep_aspect_ratio = FALSE,
  xscale = 0.9, yspace = unit(0.5, "lines"), main = NULL, sub = NULL,
  ..., residuals_type = "Pearson", gp_axis = gpar(lty = 3))
## S3 method for class 'formula'
assoc(formula, data = NULL, ..., subset = NULL, na.action = NULL, main = NULL, sub = NULL)
```

Arguments

x	a contingency table in array form with optional category labels specified in the `dimnames(x)` attribute, or an object inheriting from the `"ftable"` class (such as `"structable"` objects).
row_vars	a vector of integers giving the indices, or a character vector giving the names of the variables to be used for the rows of the association plot.
col_vars	a vector of integers giving the indices, or a character vector giving the names of the variables to be used for the columns of the association plot.
compress	logical; if FALSE, the space between the rows (columns) are chosen such that the *total* heights (widths) of the rows (columns) are all equal. If TRUE, the space between rows and columns is fixed and hence the plot is more "compressed".
xlim	a $2 \times k$ matrix of doubles, k number of total columns of the plot. The columns of xlim correspond to the columns of the association plot, the rows describe the column ranges (minimums in the first row, maximums in the second row). If xlim is NULL, the ranges are determined from the residuals according to compress (if TRUE: widest range from each column, if FALSE: from the whole association plot matrix).
ylim	a $2 \times k$ matrix of doubles, k number of total rows of the plot. The columns of ylim correspond to the rows of the association plot, the rows describe the column ranges (minimums in the first row, maximums in the second row). If ylim is NULL, the ranges are determined from the residuals according to compress (if TRUE: widest range from each row, if FALSE: from the whole association plot matrix).
spacing	a spacing object, a spacing function, or a corresponding generating function (see [strucplot](#) for more information). The default is the spacing-generating function [spacing_conditional](#) that is (by default) called with the argument list spacing_args (see spacings for more details).
spacing_args	list of arguments for the spacing-generating function, if specified (see [strucplot](#) for more information).
split_vertical	vector of logicals of length k, where k is the number of margins of x (default: FALSE). Values are recycled as needed. A TRUE component indicates that the corresponding dimension is folded into the columns, FALSE folds the dimension into the rows.

keep_aspect_ratio

                    logical indicating whether the aspect ratio should be fixed or not.

residuals_type   a character string indicating the type of residuals to be computed. Currently, only Pearson residuals are supported.

xscale   scale factor resizing the tile's width, thus adding additional space between the tiles.

yspace   object of class "unit" specifying additional space separating the rows.

gp_axis   object of class "gpar" specifying the visual aspects of the tiles' baseline.

formula   a formula object with possibly both left and right hand sides specifying the column and row variables of the flat table.

data   a data frame, list or environment containing the variables to be cross-tabulated, or an object inheriting from class table.

subset   an optional vector specifying a subset of observations to be used. Ignored if data is a contingency table.

na.action   an optional function which indicates what should happen when the data contain NAs. Ignored if data is a contingency table.

main, sub   either a logical, or a character string used for plotting the main (sub) title. If logical and TRUE, the name of the data object is used.

...   other parameters passed to strucplot

Details

Association plots have been suggested by Cohen (1980) and extended by Friendly (1992) and provide a means for visualizing the residuals of an independence model for a contingency table.

assoc is a generic function and currently has a default method and a formula interface. Both are high-level interfaces to the strucplot function, and produce (extended) association plots. Most of the functionality is described there, such as specification of the independence model, labeling, legend, spacing, shading, and other graphical parameters.

For a contingency table, the signed contribution to Pearson's χ^2 for cell $\{ij \ldots k\}$ is

$$d_{ij\ldots k} = \frac{(f_{ij\ldots k} - e_{ij\ldots k})}{\sqrt{e_{ij\ldots k}}}$$

where $f_{ij\ldots k}$ and $e_{ij\ldots k}$ are the observed and expected counts corresponding to the cell. In the association plot, each cell is represented by a rectangle that has (signed) height proportional to $d_{ij\ldots k}$ and width proportional to $\sqrt{e_{ij\ldots k}}$, so that the area of the box is proportional to the difference in observed and expected frequencies. The rectangles in each row are positioned relative to a baseline indicating independence ($d_{ij\ldots k} = 0$). If the observed frequency of a cell is greater than the expected one, the box rises above the baseline, and falls below otherwise.

Additionally, the residuals can be colored depending on a specified shading scheme (see Meyer et al., 2003). Package **vcd** offers a range of *residual-based* shadings (see the shadings help page). Some of them allow, e.g., the visualization of test statistics.

Unlike the assocplot function in the **graphics** package, this function allows the visualization of contingency tables with more than two dimensions. Similar to the construction of 'flat' tables (like objects of class "ftable" or "structable"), the dimensions are folded into rows and columns.

The layout is very flexible: the specification of shading, labeling, spacing, and legend is modularized (see strucplot for details).

Value

The "structable" visualized is returned invisibly.

Author(s)

David Meyer <David.Meyer@R-project.org>

References

Cohen, A. (1980), On the graphical display of the significant components in a two-way contingency table. *Communications in Statistics—Theory and Methods*, **A9**, 1025–1041.

Friendly, M. (1992), Graphical methods for categorical data. *SAS User Group International Conference Proceedings*, **17**, 190–200. http://datavis.ca/papers/sugi/sugi17.pdf

Meyer, D., Zeileis, A., Hornik, K. (2003), Visualizing independence using extended association plots. *Proceedings of the 3rd International Workshop on Distributed Statistical Computing*, K. Hornik, F. Leisch, A. Zeileis (eds.), ISSN 1609-395X. http://www.ci.tuwien.ac.at/Conferences/DSC-2003/Proceedings/

Meyer, D., Zeileis, A., and Hornik, K. (2006), The strucplot framework: Visualizing multi-way contingency tables with **vcd**. *Journal of Statistical Software*, **17(3)**, 1-48. URL http://www.jstatsoft.org/v17/i03/ and available as vignette("strucplot").

See Also

mosaic, strucplot, structable

Examples

```
data("HairEyeColor")
## Aggregate over sex:
(x <- margin.table(HairEyeColor, c(1, 2)))

## Ordinary assocplot:
assoc(x)
## and with residual-based shading (of independence)
assoc(x, main = "Relation between hair and eye color", shade = TRUE)

## Aggregate over Eye color:
(x <- margin.table(HairEyeColor, c(1, 3)))
chisq.test(x)
assoc(x, main = "Relation between hair color and sex", shade = TRUE)

# Visualize multi-way table
assoc(aperm(HairEyeColor), expected = ~ (Hair + Eye) * Sex,
      labeling_args = list(just_labels = c(Eye = "left"),
                           offset_labels = c(right = -0.5),
                           offset_varnames = c(right = 1.2),
```

```
                                   rot_labels = c(right = 0),
                                   tl_varnames = c(Eye = TRUE))
  )

  assoc(aperm(UCBAdmissions), expected = ~ (Admit + Gender) * Dept, compress = FALSE,
        labeling_args = list(abbreviate = c(Gender = TRUE), rot_labels = 0)
  )
```

---

assocstats   *Association Statistics*

Description

Computes the Pearson chi-Squared test, the Likelihood Ratio chi-Squared test, the phi coefficient, the contingency coefficient and Cramer's V.

Usage

```
assocstats(x)
```

Arguments

x   an $r \times c$ table.

Value

A list with components:

chisq_tests	a 2×3 table with the chi-squared statistics.
phi	The *absolute value* of the phi coefficient.
cont	The contingency coefficient.
cramer	Cramer's V.

Author(s)

David Meyer <David.Meyer@R-project.org>

References

Michael Friendly (2000), *Visualizing Categorical Data.* SAS Institute, Cary, NC.

Fleiss, J. L. (1981). *Statistical methods for rates and proportions* (2nd ed). New York: Wiley

Examples

```
data("Arthritis")
tab <- xtabs(~Improved + Treatment, data = Arthritis)
summary(assocstats(tab))
```

Baseball	*Baseball Data*

Description

Baseball data.

Usage

```
data("Baseball")
```

Format

A data frame with 322 observations and 25 variables.

name1 player's first name.

name2 player's last name.

atbat86 times at Bat: number of official plate appearances by a hitter. It counts as an official at-bat as long as the batter does not walk, sacrifice, get hit by a pitch or reach base due to catcher's interference.

hits86 hits.

homer86 home runs.

runs86 the number of runs scored by a player. A run is scored by an offensive player who advances from batter to runner and touches first, second, third and home base in that order without being put out.

rbi86 Runs Batted In: A hitter earns a run batted in when he drives in a run via a hit, walk, sacrifice (bunt or fly) fielder's choice, hit-batsman or on an error (when the official scorer rules that the run would have scored anyway).

walks86 A "walk" (or "base on balls") is an award of first base granted to a batter who receives four pitches outside the strike zone.

years Years in the Major Leagues. Seems to count all years a player has actually played in the Major Leagues, not necessarily consecutive.

atbat career times at bat.

hits career hits.

homeruns career home runs.

runs career runs.

rbi career runs batted in.

walks career walks.

league86 player's league.

div86 player's division.

team86 player's team.

posit86 player's position (see `Hitters`).

outs86 number of putouts (see `Hitters`)

assist86 number of assists (see `Hitters`)

error86 number of assists (see `Hitters`)

sal87 annual salary on opening day (in USD 1000).

league87 league in 1987.

team87 team in 1987.

Source

SAS System for Statistical Graphics, First Edition, page A2.3

References

M. Friendly (2000), *Visualizing Categorical Data*. SAS Institute, Cary, NC.

See Also

`Hitters`

Examples

```
data("Baseball")
```

BrokenMarriage *Broken Marriage Data*

Description

Data from the Danish Welfare Study about broken marriages or permanent relationships depending on gender and social rank.

Usage

```
data("BrokenMarriage")
```

Format

A data frame with 20 observations and 4 variables.

Freq frequency.

gender factor indicating gender (male, female).

rank factor indicating social rank (I, II, III, IV, V).

broken factor indicating whether the marriage or permanent relationship was broken (yes, no).

Source

E. B. Andersen (1991), The Statistical Analysis of Categorical Data, page 177.

References

E. B. Andersen (1991), *The Statistical Analysis of Categorical Data*. 2nd edition. Springer-Verlag, Berlin.

Examples

```
data("BrokenMarriage")
structable(~ ., data = BrokenMarriage)
```

Bundesliga *Ergebnisse der Fussball-Bundesliga*

Description

Results from the first German soccer league (1963-2008).

Usage

```
data("Bundesliga")
```

Format

A data frame with 14018 observations and 7 variables.

HomeTeam factor. Name of the home team.

AwayTeam factor. Name of the away team.

HomeGoals number of goals scored by the home team.

AwayGoals number of goals scored by the away team.

Round round of the game.

Year year in which the season started.

Date starting time of the game (in "POSIXct" format).

Details

The data comprises all games in the first German soccer league since its foundation in 1963. The data have been queried online from the official Web page of the DFB and prepared as a data frame in R by Daniel Dekic, Torsten Hothorn, and Achim Zeileis (replacing earlier versions of the data in the package containing only subsets of years).

Each year/season comprises 34 rounds (except 1963, 1964, 1991) so that all 18 teams play twice against each other (switching home court advantage). In 1963/64, there were only 16 teams, hence only 30 rounds. In 1991, after the German unification, there was one season with 20 teams and 38 rounds.

Source

Homepage of the Deutscher Fussball-Bund (DFB, German Football Association): http://www.dfb.de/

References

Leonhard Knorr-Held (1999), Dynamic rating of sports teams. SFB 386 "Statistical Analysis of Discrete Structures", Discussion paper **98**.

See Also

UKSoccer

Examples

```
data("Bundesliga")

## number of goals per game poisson distributed?
ngoals1 <- xtabs(~ HomeGoals, data = Bundesliga, subset = Year == 1995)
ngoals2 <- xtabs(~ AwayGoals, data = Bundesliga, subset = Year == 1995)
ngoals3 <- table(apply(subset(Bundesliga, Year == 1995)[,3:4], 1, sum))

gf1 <- goodfit(ngoals1)
gf2 <- goodfit(ngoals2)
gf3 <- goodfit(ngoals3)

summary(gf1)
summary(gf2)
summary(gf3)
plot(gf1)
plot(gf2)
plot(gf3)

Ord_plot(ngoals1)
distplot(ngoals1)
```

Bundestag2005   *Votes in German Bundestag Election 2005*

Description

Number of votes by province in the German Bundestag election 2005 (for the parties that eventually entered the parliament).

Usage

```
data("Bundestag2005")
```

Format

A 2-way `"table"` giving the number of votes for each party (`Fraktion`) in each of the 16 German provinces (`Bundesland`):

No	Name	Levels
1	Bundesland	Schleswig-Holstein, Mecklenburg-Vorpommern, ...
2	Fraktion	SPD, CDU/CSU, Gruene, FDP, Linke

Details

In the election for the German parliament "Bundestag", five parties obtained enough votes to enter the parliament: the social democrats SPD, the conservative CDU/CSU, the liberal FDP, the green party "Die Gruenen" and the leftist party "Die Linke". The table `Bundestag2005` gives the number of votes for each party (`Fraktion`) in each of the 16 German provinces (`Bundesland`). The provinces are ordered from North to South.

The data have been obtained from the German statistical office (Statistisches Bundesamt) from the Web page given below.

Note that the number of seats in the parliament cannot be computed from the number of votes alone. The examples below show the distribution of seats that resulted from the election.

Source

Der Bundeswahlleiter, Statistisches Bundesamt. <http://www.bundeswahlleiter.de/bundestagswahl2005/>

Examples

```
library(colorspace)
## The outcome of the election in terms of seats in the
## parliament was:
seats <- structure(c(226, 61, 54, 51, 222),
  .Names = c("CDU/CSU", "FDP",  "Linke", "Gruene", "SPD"))

## Hues are chosen as metaphors for the political parties
## CDU/CSU: blue, FDP: yellow, Linke: purple, Gruene: green, SPD: red
## using the respective hues from a color wheel with
## chroma = 60 and luminance = 75
parties <- rainbow_hcl(6, c = 60, l = 75)[c(5, 2, 6, 3, 1)]
names(parties) <- names(seats)
parties

## The pie chart shows that neither the SPD+Gruene coalition nor
## the opposition of CDU/CSU+FDP could assemble a majority.
## No party would enter a coalition with the leftists, leading to a
## big coalition.
pie(seats, clockwise = TRUE, col = parties)

## The regional distribution of the votes, stratified by province,
## is shown in a mosaic display: first for the 10 Western then the
## 6 Eastern provinces.
data("Bundestag2005")
votes <- Bundestag2005[c(1, 3:5, 9, 11, 13:16, 2, 6:8, 10, 12),
                        c("CDU/CSU", "FDP", "SPD", "Gruene", "Linke")]
mosaic(votes, gp = gpar(fill = parties[colnames(votes)]),
  spacing = spacing_highlighting, labeling = labeling_left,
  labeling_args = list(rot_labels = c(0, 90, 0, 0), pos_labels = "center",
  just_labels = c("center","center","center","right"), varnames = FALSE),
  margins = unit(c(2.5, 1, 1, 12), "lines"),
  keep_aspect_ratio = FALSE)
```

Butterfly   *Butterfly Species in Malaya*

Description

Data from Fisher et al. (1943) giving the number of tokens found for each of 501 species of butter-flies collected in Malaya.

Usage

```
data("Butterfly")
```

Format

A 1-way table giving the number of tokens for 501 species of butterflies. The variable and its levels are

No	Name	Levels
1	nTokens	0, 1, ..., 24

Source

Michael Friendly (2000), Visualizing Categorical Data, pages 21–22.

References

R. A. Fisher, A. S. Corbet, C. B. Williams (1943), The relation between the number of species and the number of individuals, *Journal of Animal Ecology*, **12**, 42–58.

M. Friendly (2000), *Visualizing Categorical Data*. SAS Institute, Cary, NC.

Examples

```
data("Butterfly")
Ord_plot(Butterfly)
```

---

cd_plot   *Conditional Density Plots*

Description

Computes and plots conditional densities describing how the distribution of a categorical variable y changes over a numerical variable x.

Usage

```
cd_plot(x, ...)
## Default S3 method:
cd_plot(x, y,
  plot = TRUE, ylab_tol = 0.05,
  bw = "nrd0", n = 512, from = NULL, to = NULL,
  main = "", xlab = NULL, ylab = NULL, margins = c(5.1, 4.1, 4.1, 3.1),
  gp = gpar(), name = "cd_plot", newpage = TRUE, pop = TRUE,
  ...)
## S3 method for class 'formula'
cd_plot(formula, data = list(),
  plot = TRUE, ylab_tol = 0.05,
  bw = "nrd0", n = 512, from = NULL, to = NULL,
  main = "", xlab = NULL, ylab = NULL, margins = c(5.1, 4.1, 4.1, 3.1),
  gp = gpar(), name = "cd_plot", newpage = TRUE, pop = TRUE,
  ...)
```

Arguments

x	an object, the default method expects either a single numerical variable.
y	a `"factor"` interpreted to be the dependent variable
formula	a `"formula"` of type `y ~ x` with a single dependent `"factor"` and a single numerical explanatory variable.
data	an optional data frame.
plot	logical. Should the computed conditional densities be plotted?
ylab_tol	convenience tolerance parameter for y-axis annotation. If the distance between two labels drops under this threshold, they are plotted equidistantly.
bw, n, from, to, ...	
	arguments passed to density
main, xlab, ylab	
	character strings for annotation
margins	margins when calling plotViewport
gp	a `"gpar"` object controlling the grid graphical parameters of the rectangles. It should specify in particular a vector of `fill` colors of the same length as `levels(y)`. The default is to call gray.colors.
name	name of the plotting viewport.
newpage	logical. Should grid.newpage be called before plotting?
pop	logical. Should the viewport created be popped?

Details

cd_plot computes the conditional densities of x given the levels of y weighted by the marginal distribution of y. The densities are derived cumulatively over the levels of y.

This visualization technique is similar to spinograms (see spine) but they do not discretize the explanatory variable, but rather use a smoothing approach. Furthermore, the original x axis and not

a distorted x axis (as for spinograms) is used. This typically results in conditional densities that are based on very few observations in the margins: hence, the estimates are less reliable there.

Value

The conditional density functions (cumulative over the levels of y) are returned invisibly.

Author(s)

Achim Zeileis <Achim.Zeileis@R-project.org>

References

Hofmann, H., Theus, M. (2005), *Interactive graphics for visualizing conditional distributions*, Unpublished Manuscript.

See Also

spine, density

Examples

```
## Arthritis data
data("Arthritis")
cd_plot(Improved ~ Age, data = Arthritis)
cd_plot(Improved ~ Age, data = Arthritis, bw = 3)
cd_plot(Improved ~ Age, data = Arthritis, bw = "SJ")
## compare with spinogram
spine(Improved ~ Age, data = Arthritis, breaks = 3)

## Space shuttle data
data("SpaceShuttle")
cd_plot(Fail ~ Temperature, data = SpaceShuttle, bw = 2)

## scatter plot with conditional density
cdens <- cd_plot(Fail ~ Temperature, data = SpaceShuttle, bw = 2, plot = FALSE)
plot(I(-1 * (as.numeric(Fail) - 2)) ~ jitter(Temperature, factor = 2), data = SpaceShuttle,
  xlab = "Temperature", ylab = "Failure")
lines(53:81, cdens[[1]](53:81), col = 2)
```

CoalMiners   *Breathlessness and Wheeze in Coal Miners*

Description

Data from Ashford & Sowden (1970) given by Agresti (1990) on the association between two pulmonary conditions, breathlessness and wheeze, in a large sample of coal miners who were smokers with no radiological evidence of pneumoconlosis, aged between 20–64 when examined. This data is frequently used as an example of fitting models for bivariate, binary responses.

Usage

```
data("CoalMiners")
```

Format

A 3-dimensional table of size 2 x 2 x 9 resulting from cross-tabulating variables for 18,282 coal miners. The variables and their levels are as follows:

No	Name	Levels
1	Breathlessness	B, NoB
2	Wheeze	W, NoW
3	Age	20-24, 25-29, 30-34, . . . , 60-64

Details

In an earlier version of this data set, the first group, aged 20-24, was inadvertently omitted from this data table and the breathlessness variable was called wheeze and vice versa.

Source

Michael Friendly (2000), Visualizing Categorical Data, pages 82–83, 319–322.

References

A. Agresti (1990), *Categorical Data Analysis*. Wiley-Interscience, New York, Table 7.11, p. 237

J. R. Ashford and R. D. Sowdon (1970), Multivariate probit analysis, *Biometrics*, **26**, 535–546.

M. Friendly (2000), *Visualizing Categorical Data*. SAS Institute, Cary, NC.

Examples

```
data("CoalMiners")

ftable(CoalMiners, row.vars = 3)

## Fourfold display, both margins equated
fourfold(CoalMiners[,,2:9], mfcol = c(2,4))

## Fourfold display, strata equated
fourfold(CoalMiners[,,2:9], std = "ind.max", mfcol = c(2,4))


## Log Odds Ratio Plot
summary(l <- oddsratio(CoalMiners[,,2:9]))
age <- seq(25, 60, by = 5)
plot(l,
     xlab = "Age Group",
     main = "Breathlessness and Wheeze in Coal Miners")
m <- lm(l ~ age + I(age^2))
lines(fitted(m), col = "red", lwd=2)
```

---

coindep_test   *Test for (Conditional) Independence*

Description

Performs a test of (conditional) independence of 2 margins in a contingency table by simulation from the marginal distribution of the input table under (conditional) independence.

Usage

```
coindep_test(x, margin = NULL, n = 1000,
  indepfun = function(x) max(abs(x)), aggfun = max,
  alternative = c("greater", "less"),
  pearson = TRUE)
```

Arguments

x	a contingency table.
margin	margin index(es) or corresponding name(s) of the conditioning variables. Each resulting conditional table has to be a 2-way table.
n	number of (conditional) independence tables to be drawn.
indepfun	aggregation function capturing independence in (each conditional) 2-way table.
aggfun	aggregation function aggregating the test statistics computed by indepfun.
alternative	a character string specifying the alternative hypothesis; must be either "greater" (default) or "less" (and may be abbreviated.)
pearson	logical. Should the table of Pearson residuals under independence be computed and passed to indepfun (default) or the raw table of observed frequencies?

Details

If margin is NULL this computes a simple independence statistic in a 2-way table. Alternatively, margin can give several conditioning variables and then conditional independence in the resulting conditional table is tested.

By default, this uses a (double) maximum statistic of Pearson residuals. By changing indepfun or aggfun a (maximum of) Pearson Chi-squared statistic(s) can be computed or just the usual Pearson Chi-squared statistics and so on. Other statistics can be computed by changing pearson to FALSE.

The function uses [r2dtable](#) to simulate the distribution of the test statistic under the null.

Value

A list of class "coindep_test" inheriting from "htest" with following components:

statistic	the value of the test statistic.
p.value	the p value for the test.
method	a character string indicating the type of the test.

`data.name`	a character string giving the name(s) of the data.
`observed`	observed table of frequencies
`expctd`	expected table of frequencies
`residuals`	corresponding Pearson residuals
`margin`	the `margin` used
`dist`	a vector of size n with simulated values of the distribution of the statistic under the null.
`qdist`	the corresponding quantile function (for computing critical values).
`pdist`	the corresponding distribution function (for computing p values).

Author(s)

Achim Zeileis <Achim.Zeileis@R-project.org>

See Also

[chisq.test](), [fisher.test](), [r2dtable]()

Examples

```
library(MASS)
TeaTasting <- matrix(c(3, 1, 1, 3), nr = 2,
                     dimnames = list(Guess = c("Milk", "Tea"),
                                     Truth = c("Milk", "Tea"))
)
## compute maximum statistic
coindep_test(TeaTasting)
## compute Chi-squared statistic
coindep_test(TeaTasting, indepfun = function(x) sum(x^2))
## use unconditional asymptotic distribution
chisq.test(TeaTasting, correct = FALSE)
chisq.test(TeaTasting)


data("UCBAdmissions")
## double maximum statistic
coindep_test(UCBAdmissions, margin = "Dept")
## maximum of Chi-squared statistics
coindep_test(UCBAdmissions, margin = "Dept", indepfun = function(x) sum(x^2))
## Pearson Chi-squared statistic
coindep_test(UCBAdmissions, margin = "Dept", indepfun = function(x) sum(x^2), aggfun = sum)
## use unconditional asymptotic distribution
loglm(~ Dept * (Gender + Admit), data = UCBAdmissions)
```

---

cotabplot	*Coplot for Contingency Tables*

Description

cotabplot is a generic function for creating trellis-like coplots (conditional plots) for contingency tables.

Usage

```
cotabplot(x, ...)
## Default S3 method:
cotabplot(x, cond = NULL,
  panel = cotab_mosaic, panel_args = list(),
  margins = rep(1, 4), layout = NULL,
  text_gp = gpar(fontsize = 12), rect_gp = gpar(fill = grey(0.9)),
  pop = TRUE, newpage = TRUE,
  ...)
## S3 method for class 'formula'
cotabplot(formula, data = NULL, ...)
```

Arguments

x	an object. The default method can deal with contingency tables in array form.
cond	margin index(es) or corresponding name(s) of the conditioning variables.
panel	panel function applied for each conditioned plot, see details.
panel_args	list of arguments passed to panel if this is a panel-generating function inheriting from class "grapcon_generator".
margins	either an object of class "unit" of length 4, or a numeric vector of length 4. The elements are recycled as needed. giving the margins around the whole plot.
layout	integer vector (of length two), giving the number of rows and columns for the panel.
text_gp	object of class "gpar" used for the text in the panel titles.
rect_gp	object of class "gpar" used for the rectangles with the panel titles.
pop	logical indicating whether the generated viewport tree should be removed at the end of the drawing or not.
newpage	logical controlling whether a new grid page should be created.
...	further arguments passed to the panel-generating function.
formula	a formula specifying the variables used to create a contingency table from data. It has to be of type ~ x + y | z where z is/are the conditioning variable(s) used.
data	either a data frame, or an object of class "table" or "ftable".

Details

cotabplot is a generic function designed to create coplots or conditional plots (see Cleveland, 1993, and Becker, Cleveland, Shyu, 1996) similar to coplot but for contingency tables.

cotabplot takes on computing the conditioning information and setting up the trellis display, and then relies on a panel function to create plots from the full table and the conditioning information. A simple example would be a contingency table tab with margin names "x", "y" and "z". To produce this plot either the default interface can be used or the formula interface via

cotabplot(tab, "z") cotabplot(~ x + y | z, data = tab)

The panel function needs to be of the form

panel(x, condlevels)

where x is the *full* table (tab in the example above) and condlevels is a named vector with the levels (e.g., c(z = "z1") in the example above).

Alternatively, panel can also be a panel-generating function of class "grapcon_generator" which creates a function with the interface described above. The panel-generating function is called with the interface

panel(x, condvars, ...)

where again x is the full table, condvars is now only a vector with the names of the conditioning variables (and not their levels, e.g., "z" in the example above). Further arguments can be passed to the panel-generating function via ... which also includes the arguments set in panel_args.

Suitable panel-generating functions for mosaic, association and sieve plots can be found at cotab_mosaic.

A description of the underlying ideas is given in Zeileis, Meyer, Hornik (2005).

Author(s)

Achim Zeileis <Achim.Zeileis@R-project.org>

References

Becker, R.A., Cleveland, W.S., Shyu, M.-J. (1996), The visual design and control of trellis display. *Journal of Computational and Graphical Statistics*, **5**, 123–155.

Cleveland, W.S. (1993), *Visualizing Data*, Summit, New Jersey: Hobart Press.

Meyer, D., Zeileis, A., and Hornik, K. (2006), The strucplot framework: Visualizing multi-way contingency tables with **vcd**. *Journal of Statistical Software*, **17(3)**, 1-48. URL http://www.jstatsoft.org/v17/i03/ and available as vignette("strucplot").

Zeileis, A., Meyer, D., Hornik K. (2007), *Residual-based shadings for visualizing (conditional) independence*, *Journal of Computational and Graphical Statistics*, **16**, 507–525.

See Also

cotab_mosaic, cotab_coindep, co_table, coindep_test

Examples

```
data("UCBAdmissions")

cotabplot(~ Admit + Gender | Dept, data = UCBAdmissions)
cotabplot(~ Admit + Gender | Dept, data = UCBAdmissions, panel = cotab_assoc)

ucb <- cotab_coindep(UCBAdmissions, condvars = "Dept", type = "assoc",
                     n = 5000, margins = c(3, 1, 1, 3))
cotabplot(~ Admit + Gender | Dept, data = UCBAdmissions, panel = ucb)
```

---

cotab_panel   *Panel-generating Functions for Contingency Table Coplots*

Description

Panel-generating functions visualizing contingency tables that can be passed to `cotabplot`.

Usage

```
cotab_mosaic(x = NULL, condvars = NULL, ...)
cotab_assoc(x = NULL, condvars = NULL, ylim = NULL, ...)
cotab_sieve(x = NULL, condvars = NULL, ...)
cotab_fourfold(x = NULL, condvars = NULL, ...)
cotab_coindep(x, condvars,
  test = c("doublemax", "maxchisq", "sumchisq"),
  level = NULL, n = 1000, interpolate = c(2, 4),
  h = NULL, c = NULL, l = NULL, lty = 1,
  type = c("mosaic", "assoc"), legend = FALSE, ylim = NULL, ...)
```

Arguments

x	a contingency tables in array form.
condvars	margin name(s) of the conditioning variables.
ylim	y-axis limits for `assoc` plot. By default this is computed from `x`.
test	character indicating which type of statistic should be used for assessing conditional independence.
level,n,h,c,l,lty,interpolate	
	variables controlling the HCL shading of the residuals, see [shadings](#) for more details.
type	character indicating which type of plot should be produced.
legend	logical. Should a legend be produced in each panel?
...	further arguments passed to the plotting function (such as [mosaic](#) or [assoc](#) or [sieve](#) respectively).

Details

These functions of class "panel_generator" are panel-generating functions for use with cotabplot, i.e., they return functions with the interface

panel(x, condlevels)

required for cotabplot. The functions produced by cotab_mosaic, cotab_assoc and cotab_sieve essentially only call co_table to produce the conditioned table and then call mosaic, assoc or sieve respectively with the arguments specified.

The function cotab_coindep is similar but additionally chooses an appropriate residual-based shading visualizing the associated conditional independence model. The conditional independence test is carried out via coindep_test and the shading is set up via shading_hcl.

A description of the underlying ideas is given in Zeileis, Meyer, Hornik (2005).

Author(s)

Achim Zeileis <Achim.Zeileis@R-project.org>

References

Meyer, D., Zeileis, A., and Hornik, K. (2006), The strucplot framework: Visualizing multi-way contingency tables with **vcd**. *Journal of Statistical Software*, **17(3)**, 1-48. URL http://www.jstatsoft.org/v17/i03/ and available as vignette("strucplot").

Zeileis, A., Meyer, D., Hornik K. (2007), *Residual-based shadings for visualizing (conditional) independence*, *Journal of Computational and Graphical Statistics*, **16**, 507–525.

See Also

cotabplot, mosaic, assoc, sieve, co_table, coindep_test, shading_hcl

Examples

```
data("UCBAdmissions")

cotabplot(~ Admit + Gender | Dept, data = UCBAdmissions)
cotabplot(~ Admit + Gender | Dept, data = UCBAdmissions, panel = cotab_assoc)
cotabplot(~ Admit + Gender | Dept, data = UCBAdmissions, panel = cotab_fourfold)

ucb <- cotab_coindep(UCBAdmissions, condvars = "Dept", type = "assoc",
                     n = 5000, margins = c(3, 1, 1, 3))
cotabplot(~ Admit + Gender | Dept, data = UCBAdmissions, panel = ucb)
```

---

co_table   *Compute Conditional Tables*

Description

For a contingency table in array form, compute a list of conditional tables given some margins.

Usage

```
co_table(x, margin, collapse = ".")
```

Arguments

x	a contingency table in array form.
margin	margin index(es) or corresponding name(s) of the conditioning variables.
collapse	character used when collapsing level names (if more than 1 margin is specified).

Details

This is essentially an interface to [which is more convenient for arrays of arbitrary dimension.

Value

A list of the resulting conditional tables.

Author(s)

Achim Zeileis <Achim.Zeileis@R-project.org>

Examples

```
data("HairEyeColor")
co_table(HairEyeColor, 1)
co_table(HairEyeColor, c("Hair", "Eye"))
co_table(HairEyeColor, 1:2, collapse = "")
```

DanishWelfare   *Danish Welfare Study Data*

Description

Data from the Danish Welfare Study.

Usage

```
data("DanishWelfare")
```

Format

A data frame with 180 observations and 5 variables.

Freq frequency.

Alcohol factor indicating daily alcohol consumption: less than 1 unit (<1), 1-2 units (1-2) or more
than 2 units (>2). 1 unit is approximately 1 bottle of beer or 4cl 40% alcohol.

Income factor indicating income group in 1000 DKK (0-50, 50-100, 100-150, >150).

Status factor indicating marriage status (Widow, Married, Unmarried).

Urban factor indicating urbanization: Copenhagen (Copenhagen), Suburbian Copenhagen (Sub-
Copenhagen), three largest cities (LargeCity), other cities (City), countryside (Country).

Source

E. B. Andersen (1991), The Statistical Analysis of Categorical Data, page 205.

References

E. B. Andersen (1991), *The Statistical Analysis of Categorical Data*. 2nd edition. Springer-Verlag,
Berlin.

Examples

```
data("DanishWelfare")
ftable(xtabs(Freq ~ ., data = DanishWelfare))
```

---

distplot   *Diagnostic Distribution Plots*

Description

Diagnostic distribution plots: poissonness, binomialness and negative binomialness plots.

Usage

```
distplot(x, type = c("poisson", "binomial", "nbinomial"),
  size = NULL, lambda = NULL, legend = TRUE, xlim = NULL, ylim = NULL,
  conf_int = TRUE, conf_level = 0.95, main = NULL,
  xlab = "Number of occurrences", ylab = "Distribution metameter",
  gp = gpar(cex = 0.8), lwd=2, name = "distplot", newpage = TRUE, pop = TRUE, ...)
```

Arguments

x	either a vector of counts, a 1-way table of frequencies of counts or a data frame or matrix with frequencies in the first column and the corresponding counts in the second column.
type	a character string indicating the distribution.
size	the size argument for the binomial and negative binomial distribution. If set to NULL and type is "binomial", then size is taken to be the maximum count. If set to NULL and type is "nbinomial", then size is estimated from the data.
lambda	parameter of the poisson distribution. If type is "poisson" and lambda is specified a leveled poissonness plot is produced.
legend	logical. Should a legend be plotted?
xlim	limits for the x axis.
ylim	limits for the y axis.
conf_int	logical. Should confidence intervals be plotted?
conf_level	confidence level for confidence intervals.
main	a title for the plot.
xlab	a label for the x axis.
ylab	a label for the y axis.
gp	a "gpar" object controlling the grid graphical parameters of the points.
lwd	line width for the fitted line
name	name of the plotting viewport.
newpage	logical. Should [grid.newpage](#) be called before plotting?
pop	logical. Should the viewport created be popped?
...	further arguments passed to [grid.points](#).

Details

distplot plots the number of occurrences (counts) against the distribution metameter of the specified distribution. If the distribution fits the data, the plot should show a straight line. See Friendly (2000) for details.

In these plots, the open points show the observed count metameters; the filled points show the confidence interval centers, and the dashed lines show the conf_level confidence intervals for each point.

Value

Returns invisibly a data frame containing the counts (Counts), frequencies (Freq) and other details of the computations used to construct the plot.

Author(s)

Achim Zeileis <Achim.Zeileis@R-project.org>

References

D. C. Hoaglin (1980), A poissonness plot, *The American Statistican*, **34**, 146–149.

D. C. Hoaglin & J. W. Tukey (1985), Checking the shape of discrete distributions. In D. C. Hoaglin, F. Mosteller, J. W. Tukey (eds.), *Exploring Data Tables, Trends and Shapes*, chapter 9. John Wiley & Sons, New York.

M. Friendly (2000), *Visualizing Categorical Data*. SAS Institute, Cary, NC.

Examples

```
## Simulated data examples:
dummy <- rnbinom(1000, size = 1.5, prob = 0.8)
distplot(dummy, type = "nbinomial")

## Real data examples:
data("HorseKicks")
data("Federalist")
data("Saxony")
distplot(HorseKicks, type = "poisson")
distplot(HorseKicks, type = "poisson", lambda = 0.61)
distplot(Federalist, type = "poisson")
distplot(Federalist, type = "nbinomial", size = 1)
distplot(Federalist, type = "nbinomial")
distplot(Saxony, type = "binomial", size = 12)
```

---

doubledecker   *Doubledecker Plot*

Description

This function creates a doubledecker plot visualizing a classification rule.

Usage

```
## S3 method for class 'formula'
doubledecker(formula, data = NULL, ..., main = NULL)
## Default S3 method:
doubledecker(x, depvar = length(dim(x)),
  margins = c(1,4, length(dim(x)) + 1, 1),
  gp = gpar(fill = rev(gray.colors(tail(dim(x), 1)))),
  labeling = labeling_doubledecker,
  spacing = spacing_highlighting,
  main = NULL, keep_aspect_ratio = FALSE, ...)
```

Arguments

formula	a formula specifying the variables used to create a contingency table from `data`. The dependent variable is used last for splitting.
data	either a data frame, or an object of class `"table"` or `"ftable"`.
x	a contingency table in array form, with optional category labels specified in the `dimnames(x)` attribute.
depvar	dimension index or character string specifying the dependent variable. That will be sorted last in the table.
margins	margins of the plot. Note that by default, all factor names (except the last one) and their levels are visualized *as a block* under the plot.
gp	object of class `"gpar"` used for the tiles of the last variable.
labeling	labeling function or corresponding generating generating function (see [strucplot](#) for details).
spacing	spacing object, spacing function or corresponding generating function (see [strucplot](#) for details).
main	either a logical, or a character string used for plotting the main title. If `main` is `TRUE`, the name of the `data` object is used.
keep_aspect_ratio	
	logical indicating whether the aspect ratio should be maintained or not.
...	Further parameters passed to `mosaic`.

Details

Doubledecker plots visualize the the dependence of one categorical (typically binary) variable on further categorical variables. Formally, they are mosaic plots with vertical splits for all dimensions (antecedents) except the last one, which represents the dependent variable (consequent). The last variable is visualized by horizontal splits, no space between the tiles, and separate colors for the levels.

Value

The `"structable"` visualized is returned invisibly.

Author(s)

David Meyer <David.Meyer@R-project.org>

References

H. Hoffmann (2001), Generalized odds ratios for visual modeling. *Journal of Computational and Graphical Statistics*, **10**, 4, 628–640.

Meyer, D., Zeileis, A., and Hornik, K. (2006), The strucplot framework: Visualizing multi-way contingency tables with **vcd**. *Journal of Statistical Software*, **17(3)**, 1-48. URL http://www.jstatsoft.org/v17/i03/ and available as vignette(`"strucplot"`).

See Also

strucplot, mosaic

Examples

```
data("Titanic")
doubledecker(Titanic)
doubledecker(Titanic, depvar = "Survived")
doubledecker(Survived ~ ., data = Titanic)
```

Employment	*Employment Status*

Description

Data from a 1974 Danish study given by Andersen (1991) on the employees who had been laid off. The workers are classified by their employment status on 1975-01-01, the cause of their layoff and the length of employment before they were laid off.

Usage

```
data("Employment")
```

Format

A 3-dimensional array resulting from cross-tabulating variables for 1314 employees. The variables and their levels are as follows:

No	Name	Levels
1	EmploymentStatus	NewJob, Unemployed
2	EmploymentLength	<1Mo, 1-3Mo, 3-12Mo, 1-2Yr, 2-5Yr, >5Yr
3	LayoffCause	Closure, Replaced

Source

Michael Friendly (2000), Visualizing Categorical Data, pages 126–129.

References

E. B. Andersen (1991), *The Statistical Analysis of Categorical Data*. Springer-Verlag, Berlin.

M. Friendly (2000), *Visualizing Categorical Data*. SAS Institute, Cary, NC.

Examples

```
data("Employment")

## Employment Status
mosaic(Employment,
       expected = ~ LayoffCause * EmploymentLength + EmploymentStatus,
       main = "Layoff*EmployLength + EmployStatus")

mosaic(Employment,
       expected = ~ LayoffCause * EmploymentLength + LayoffCause * EmploymentStatus,
       main = "Layoff*EmployLength + Layoff*EmployStatus")

## Stratified view

grid.newpage()
pushViewport(viewport(layout = grid.layout(ncol = 2)))
pushViewport(viewport(layout.pos.col = 1))

## Closure
mosaic(Employment[,,1], main = "Layoff: Closure", newpage = FALSE)

popViewport(1)
pushViewport(viewport(layout.pos.col = 2))

## Replaced
mosaic(Employment[,,2], main = "Layoff: Replaced", newpage = FALSE)
popViewport(2)
```

Federalist   *'May' in Federalist Papers*

Description

Data from Mosteller & Wallace (1984) investigating the use of certain keywords ('may' in this data set) to identify the author of 12 disputed 'Federalist Papers' by Alexander Hamilton, John Jay and James Madison.

Usage

```
data("Federalist")
```

Format

A 1-way table giving the number of occurrences of 'may' in 262 blocks of text. The variable and its levels are

No	Name	Levels
1	nMay	$0, 1, \ldots, 6$

Source

Michael Friendly (2000), Visualizing Categorical Data, page 19.

References

F. Mosteller & D. L. Wallace (1984), *Applied Bayesian and Classical Inference: The Case of the Federalist Papers*. Springer-Verlag, New York, NY.

M. Friendly (2000), *Visualizing Categorical Data*. SAS Institute, Cary, NC.

Examples

```
data("Federalist")
gf <- goodfit(Federalist, type = "nbinomial")
summary(gf)
plot(gf)
```

---

fourfold   *Fourfold Plots*

Description

Creates an (extended) fourfold display of a $2 \times 2 \times k$ contingency table, allowing for the visual inspection of the association between two dichotomous variables in one or several populations (strata).

Usage

```
fourfold(x,
  color = c("#99CCFF", "#6699CC", "#FFA0A0", "#A0A0FF", "#FF0000", "#000080"),
  conf_level = 0.95, std = c("margins", "ind.max", "all.max"),
  margin = c(1, 2), space = 0.2, main = NULL, sub = NULL,
  mfrow = NULL, mfcol = NULL, extended = TRUE, ticks = 0.15,
  p_adjust_method = p.adjust.methods, newpage = TRUE,
  fontsize = 12, default_prefix = c("Row", "Col", "Strata"), sep = ": ")
```

Arguments

x	a $2 \times 2 \times k$ contingency table in array form, or a 2×2 matrix if k is 1. If `length(dim(x)>3`, dimensions `3:length(dim(x)` are silently raveled into a combined strata dimension with k=prod(dim(x)[-(1:2)]).
color	a vector of length 6 specifying the colors to use for the smaller and larger diagonals of each 2×2 table. The first pair is used for the standard (non-extended) plots, the other two for the extended version: the second/third pair is used for tables with non-significant/significant log-odds ratios, respectively, the latter being visualized in brighter colors.
conf_level	confidence level used for the confidence rings on the odds ratios. Must be a single non-negative number less than 1; if set to 0, confidence rings are suppressed.
std	a character string specifying how to standardize the table. Must be one of `"margins"`, `"ind.max"`, or `"all.max"`, and can be abbreviated by the initial letter. If set to `"margins"`, each 2×2 table is standardized to equate the margins specified by margin while preserving the odds ratio. If `"ind.max"` or `"all.max"`, the tables are either individually or simultaneously standardized to a maximal cell frequency of 1.
margin	a numeric vector with the margins to equate. Must be one of 1, 2, or `c(1, 2)` (the default), which corresponds to standardizing only the row, only column, or both row and column in each 2×2 table. Only used if std equals `"margins"`.
space	the amount of space (as a fraction of the maximal radius of the quarter circles) used for the row and column labels.
main, sub	character string for the fourfold plot title/subtitle.
mfrow, mfcol	a numeric vector with two components: *nr* and *nc*, indicating that the displays for the 2×2 tables should be arranged in an *nr* by *nc* layout, filled by rows/columns. The defaults are calculated to give a collection of plots in landscape orientation when *k* is not a perfect square.
extended	logical; if TRUE, extended plots are plotted, i.e., colors are brighter for significant log-odds ratios, and ticks are plotted showing the direction of association for positive log-odds.
ticks	the length of the ticks. If set to 0, no ticks are plotted.
p_adjust_method	
	method to be used for p-value adjustments for multi-stratum plots, as provided by link[stats]{p.adjust}. Use p_adjust_method="none" to disable this adjustment. The p-values are used for the 'visual' significance tests of the odds ratios.

newpage	logical; if TRUE, grid.newpage() is called before plotting.
fontsize	fontsize of main title. Other labels are scaled relative to this.
default_prefix	character vector of length 3 with default labels for possibly missing row/column/strata variable names.
sep	default separator between variable names and levels for labels.

Details

The fourfold display is designed for the display of $2 \times 2 \times k$ tables.

Following suitable standardization, the cell frequencies f_{ij} of each 2×2 table are shown as a quarter circle whose radius is proportional to $\sqrt{f_{ij}}$ so that its area is proportional to the cell frequency. An association (odds ratio different from 1) between the binary row and column variables is indicated by the tendency of diagonally opposite cells in one direction to differ in size from those in the other direction; color is used to show this direction. Confidence rings for the odds ratio allow a visual test of the null of no association; the rings for adjacent quadrants overlap iff the observed counts are consistent with the null hypothesis.

Typically, the number k corresponds to the number of levels of a stratifying variable, and it is of interest to see whether the association is homogeneous across strata. The fourfold display visualizes the pattern of association. Note that the confidence rings for the individual odds ratios are not adjusted for multiple testing.

References

Friendly, M. (1994), *A fourfold display for 2 by 2 by k tables*. Technical Report 217, York University, Psychology Department, http://datavis.ca/papers/4fold/4fold.pdf.

Friendly, M. (2000), *Visualizing Categorical Data*. SAS Institute, Cary, NC.

See Also

mosaic, assoc

link[stats]{p.adjust} for methods of p value adjustment

Examples

```
data("UCBAdmissions")
## Use the Berkeley admission data as in Friendly (1995).
x <- aperm(UCBAdmissions, c(2, 1, 3))
dimnames(x)[[2]] <- c("Yes", "No")
names(dimnames(x)) <- c("Sex", "Admit?", "Department")
ftable(x)

## Fourfold display of data aggregated over departments, with
## frequencies standardized to equate the margins for admission
## and sex.
## Figure 1 in Friendly (1994).
fourfold(margin.table(x, c(1, 2)))

## Fourfold display of x, with frequencies in each table
```

```
## standardized to equate the margins for admission and sex.
## Figure 2 in Friendly (1994).
fourfold(x)
cotabplot(x, panel = cotab_fourfold)

## Fourfold display of x, with frequencies in each table
## standardized to equate the margins for admission. but not
## for sex.
## Figure 3 in Friendly (1994).
fourfold(x, margin = 2)
```

---

goodfit   *Goodness-of-fit Tests for Discrete Data*

Description

Fits a discrete (count data) distribution for goodness-of-fit tests.

Usage

```
goodfit(x, type = c("poisson", "binomial", "nbinomial"),
  method = c("ML", "MinChisq"), par = NULL)
## S3 method for class 'goodfit'
predict(object, newcount = NULL, type = c("response", "prob"), ...)
```

Arguments

x	either a vector of counts, a 1-way table of frequencies of counts or a data frame or matrix with frequencies in the first column and the corresponding counts in the second column.
type	a character string indicating which distribution should be fit (for goodfit) or indicating the type of prediction (fitted response or probabilities in predict) respectively.
method	a character string indicating whether the distribution should be fit via ML (Maximum Likelihood) or Minimum Chi-squared.
par	a named list giving the distribution parameters (named as in the corresponding density function), if set to NULL, the default, the parameters are estimated. If the parameter size is not specified if type is "binomial" it is taken to be the maximum count. If type is "nbinomial", then parameter size can be specified to fix it so that only the parameter prob will be estimated (see the examples below).
object	an object of class "goodfit".
newcount	a vector of counts. By default the counts stored in object are used, i.e., the fitted values are computed. These can also be extracted by fitted(object).
...	*currently not used*.

Details

goodfit essentially computes the fitted values of a discrete distribution (either Poisson, binomial or negative binomial) to the count data given in x. If the parameters are not specified they are estimated either by ML or Minimum Chi-squared.

To fix parameters, par should be a named list specifying the parameters lambda for "poisson" and prob and size for "binomial" or "nbinomial", respectively. If for "binomial", size is not specified it is not estimated but taken as the maximum count.

The corresponding Pearson Chi-squared or likelihood ratio statistic, respectively, is computed and given with their p values by the summary method. The summary method always prints this information and returns a matrix with the printed information invisibly. The plot method produces a rootogram of the observed and fitted values.

In case of count distribtions (Poisson and negative binomial), the minimum Chi-squared approach is somewhat ad hoc. Strictly speaking, the Chi-squared asymptotics would only hold if the number of cells were fixed or did not increase too quickly with the sample size. However, in goodfit the number of cells is data-driven: Each count is a cell of its own. All counts larger than the maximal count are merged into the cell with the last count for computing the test statistic.

Value

A list of class "goodfit" with elements:

observed	observed frequencies.
count	corresponding counts.
fitted	expected frequencies (fitted by ML).
type	a character string indicating the distribution fitted.
method	a character string indicating the fitting method (can be either "ML", "MinChisq" or "fixed" if the parameters were specified).
df	degrees of freedom.
par	a named list of the (estimated) distribution parameters.

Author(s)

Achim Zeileis <Achim.Zeileis@R-project.org>

References

M. Friendly (2000), *Visualizing Categorical Data*. SAS Institute, Cary, NC.

See Also

rootogram

Examples

```
## Simulated data examples:
dummy <- rnbinom(200, size = 1.5, prob = 0.8)
gf <- goodfit(dummy, type = "nbinomial", method = "MinChisq")
summary(gf)
plot(gf)

dummy <- rbinom(100, size = 6, prob = 0.5)
gf1 <- goodfit(dummy, type = "binomial", par = list(size = 6))
gf2 <- goodfit(dummy, type = "binomial", par = list(prob = 0.6, size = 6))
summary(gf1)
plot(gf1)
summary(gf2)
plot(gf2)

## Real data examples:
data("HorseKicks")
HK.fit <- goodfit(HorseKicks)
summary(HK.fit)
plot(HK.fit)

data("Federalist")
## try geometric and full negative binomial distribution
F.fit <- goodfit(Federalist, type = "nbinomial", par = list(size = 1))
F.fit2 <- goodfit(Federalist, type = "nbinomial")
summary(F.fit)
summary(F.fit2)
plot(F.fit)
plot(F.fit2)
```

---

grid_barplot   *Barplot*

Description

Bar plots of 1-way tables in grid.

Usage

```
grid_barplot(height, width = 0.8, offset = 0,
  names = NULL, xlim = NULL, ylim = NULL, xlab = "", ylab = "", main = "",
  gp = gpar(fill = "lightgray"), name = "grid_barplot",
  newpage = TRUE, pop = FALSE)
```

Arguments

height	either a vector or a 1-way table of frequencies.
width	width of the bars (recycled if needed to the number of bars).

offset	offset of the bars (recycled if needed to the number of bars).
names	a vector of names for the bars, if set to NULL the names of `height` are used.
xlim	limits for the x axis.
ylim	limits for the y axis.
xlab	a label for the x axis.
ylab	a label for the y axis.
main	a title for the plot.
gp	a "gpar" object controlling the grid graphical parameters of the rectangles.
name	name of the plotting viewport.
newpage	logical. Should `grid.newpage` be called before plotting?
pop	logical. Should the viewport created be popped?

Details

`grid_barplot` mimics (some of) the features of `barplot`, but currently it only supports 1-way tables.

Author(s)

Achim Zeileis <Achim.Zeileis@R-project.org>

Examples

```
grid_barplot(sample(1:6), names = letters[1:6])
```

---

grid_legend	*Legend Function for grid Graphics*

Description

This function can be used to add legends to *grid-based* plots.

Usage

```
grid_legend(x, y, pch = NA, col = par('col'), labels, frame = TRUE,
            hgap = unit(0.8, "lines"), vgap = unit(0.8, "lines"),
            default_units = "lines", gp = gpar(), draw = TRUE,
            title = NULL, just = 'center', lwd = NA, lty = NA,
            gp_title = NULL, gp_labels = NULL,
            gp_frame = gpar(fill = "transparent"))
```

Arguments

x	character string 'topright', 'topleft', 'bottomright', 'bottomleft' or x coordinate of the legend
y	y coordinates of the legend
pch	integer vector of plotting symbols if any
col	character vector of colors for the symbols
labels	character vector of labels corresponding to the symbols
frame	logical indicating whether the legend should have a border or not.
hgap	object of class "unit" specifying the space between symbols and labels
vgap	object of class "unit" specifying the space between the lines
default_units	character string indicating the default unit
gp	object of class "gpar" used for the legend
draw	logical indicating whether the legend be drawn or not
title	character string indicating the plot's title
just	justification of the legend relative to its (x, y) location. see ?viewport for more details
lwd	positive number to set the line width. if specified lines are drawn
lty	line type. if specified lines are drawn
gp_title	object of class "gpar" used for the title
gp_labels	object of class "gpar" used for the labels
gp_frame	object of class "gpar" used for the frame

Value

Invisibly, the legend as a "grob" object.

Author(s)

David Meyer <David.Meyer@R-project.org> Florian Gerber <florian.gerber@math.uzh.ch>

See Also

[legend](#)

Examples

```
data("Lifeboats")
attach(Lifeboats)
ternaryplot(Lifeboats[,4:6],
  pch = ifelse(side == "Port", 1, 19),
  col = ifelse(side == "Port", "red", "blue"),
  id  = ifelse(men / total > 0.1, as.character(boat), NA),
  prop_size = 2,
  dimnames_position = "edge",
```

```
      main = "Lifeboats on Titanic")
grid_legend(0.8, 0.9, c(1, 19), c("red", "blue"),
  c("Port", "Starboard"), title = "SIDE")


grid.newpage()
pushViewport(viewport(height = .9, width = .9 ))
grid.rect(gp = gpar(lwd = 2, lty = 2))

grid_legend(x = unit(.05,'npc'),
            y = unit(.05,'npc'),
            just = c(0,0),
            pch = c(1,2,3),
            col = c(1,2,3),
            lwd=NA,
            lty=NA,
            labels = c("b",'r','g'),
            title = NULL,
            gp=gpar(lwd=2, cex=1),
            hgap = unit(.8, "lines"),
            vgap = unit(.9, "lines"))

grid_legend(x = unit(1,'npc'),
            y = unit(1,'npc'),
            just = c(1,1),
            pch = NA,
            col = c(1,2,3,4),
            lwd=c(1,1,1,3),
            lty=c(1,2,1,3),
            labels = c("black",'red','green','blue'),
          gp_labels = list(gpar(col = 1), gpar(col = 2), gpar(col = 3), gpar(col = 4)),
            title = NULL,
            gp=gpar(lwd=2, cex=1),
            hgap = unit(.8, "lines"),
            vgap = unit(.9, "lines"))

grid_legend(x = 'topleft',
            pch = c(1,NA,2,NA),
            col = c(1,2,3,4),
            lwd=NA,
            lty=c(NA,2,NA,3),
            labels = c("black",'red','green','blue'),
            title = 'Some LONG Title',
            gp_title = gpar(col = 3),
            gp_frame = gpar(col = 4, lty = 2, fill = "transparent"),
            gp_labels = gpar(col = 6),
            gp=gpar(lwd=2, cex=2, col = 1),
            hgap = unit(.8, "lines"),
            vgap = unit(.9, "lines"))

grid_legend(x = .7,
            y = .7,
            pch = c(1,NA,2,NA),
```

```
               col = c(1,2,3,4),
               lwd=1,
               lty=c(NA,2,NA,3),
               labels = c("black",'red','green','blue'),
               title = 'short T',
               gp=gpar(lwd=1, cex=.7,col = 1),
               hgap = unit(.8, "lines"),
               vgap = unit(.9, "lines"))

   grid_legend(x = 'bottomright',
               pch = c(1,NA,2,NA),
               col = c(2),
               lwd=NA,
               lty=c(NA,2,NA,3),
               labels = c("black",'red','green','blue'),
               title = NULL,
               gp=gpar(lwd=2, cex=1,col = 1),
               hgap = unit(.8, "lines"),
               vgap = unit(.9, "lines"))
```

Hitters   *Hitters Data*

Description

This data set is deduced from the [Baseball](#) fielding data set: fielding performance basically includes the numbers of Errors, Putouts and Assists made by each player. In order to reduce the number of observations, the was compressed by calculating the mean number of errors, putouts and assists for each team and for only 6 positions (1B, 2B, 3B, C, OF, SS and UT). In addition, each of these three variables was scaled to a common range by dividing each variable by the maximum of the variable.

Usage

```
data("Hitters")
```

Format

A data frame with 154 observations and 4 variables.

Positions factor indicating the field position (1B=first baseman, 2B=second baseman, 3B=third baseman, C=catcher, OF=outfielder, SS=Short Stop, UT=Utility Players).

Putouts occur when a fielder causes an opposing player to be tagged or forced out.

Assists are credited to other fielders involved in making that putout.

Errors count the errors made by a player.

Source

SAS System for Statistical Graphics, First Edition, Page A2.3

References

M. Friendly (2000), *Visualizing Categorical Data*. SAS Institute, Cary, NC.

Examples

```
data("Hitters")
attach(Hitters)

colors <- c("black","red","green","blue","red","black","blue")
pch <- substr(levels(Positions), 1, 1)
ternaryplot(Hitters[,2:4],
  pch = as.character(Positions),
  col = colors[as.numeric(Positions)],
  main = "Baseball Hitters Data")
grid_legend(0.8, 0.9, pch, colors, levels(Positions),
  title = "POSITION(S)")

detach(Hitters)
```

---

hls   *HLS Color Specification*

Description

Create a HLS color from specifying hue, luminance and saturation.

Usage

```
hls(h = 1, l = 0.5, s = 1)
```

Arguments

h	hue value in [0, 1].
l	luminance value in [0, 1].
s	saturation value in [0, 1].

Details

HLS colors are a similar specification of colors as HSV colors, but using hue/luminance/saturation rather that hue/saturation/value.

Author(s)

Achim Zeileis <Achim.Zeileis@R-project.org>

See Also

hsv, hcl2hex, polarLUV

Examples

```
## an HLS color wheel
pie(rep(1, 12), col = sapply(1:12/12, function(x) hls(x)))
```

HorseKicks   *Death by Horse Kicks*

Description

Data from von Bortkiewicz (1898), given by Andrews \& Herzberg (1985), on number of deaths by horse or mule kicks in 10 (of 14 reported) corps of the Prussian army. 4 corps were not considered by Fisher (1925) as they had a different organization. This data set is a popular subset of the [VonBort](#) data.

Usage

```
data("HorseKicks")
```

Format

A 1-way table giving the number of deaths in 200 corps-years. The variable and its levels are

No	Name	Levels
1	nDeaths	$0, 1, \ldots, 4$

Source

Michael Friendly (2000), Visualizing Categorical Data, page 18.

References

D. F. Andrews & A. M. Herzberg (1985), *Data: A Collection of Problems from Many Fields for the Student and Research Worker*. Springer-Verlag, New York, NY.

R. A. Fisher (1925), *Statistical Methods for Research Workers*. Oliver \& Boyd, London.

L. von Bortkiewicz (1898), *Das Gesetz der kleinen Zahlen*. Teubner, Leipzig.

M. Friendly (2000), *Visualizing Categorical Data*. SAS Institute, Cary, NC.

See Also

[VonBort](#)

Examples

```
data("HorseKicks")
gf <- goodfit(HorseKicks)
summary(gf)
```

```
plot(gf)
```

```
Hospital                    Hospital data
```

Description

The table relates the length of stay (in years) of 132 long-term schizophrenic patients in two London mental hospitals with the frequency of visits.

Usage

```
data("Hospital")
```

Format

A 2-dimensional array resulting from cross-tabulating 132 patients. The variables and their levels are as follows:

No	Name	Levels
1	Visit Frequency	Regular, Less than monthly, Never
2	Length of Stay	2–9 years, 10–19 years, 20+ years

Details

Wing (1962) who collected this data concludes that the longer the length of stay in hospital, the less frequent the visits.

Haberman (1974) notes that this pattern does not increase from the "Less than monthly" to the "Never" group, which are homogeneous.

Source

S.J Haberman (1974): Log-linear models for frequency tables with ordered classifications. Biometrics, 30:689–700.

References

J.K. Wing (1962): Institutionalism in mental hospitals. British Journal of Social Clinical Psychology, 1:38–51.

Examples

```
data("Hospital")

mosaic(t(Hospital), shade = TRUE)
mosaic(Hospital, shade = TRUE)
sieve(Hospital, shade = TRUE)
assoc(Hospital, shade = TRUE)
```

---

independence_table *Independence Table*

Description

Computes table of expected frequencies (under the null hypotheses of independence) from an n-way table.

Usage

```
independence_table(x, frequency = c("absolute", "relative"))
```

Arguments

x   a table.

frequency   indicates whether absolute or relative frequencies should be computed.

Value

A table with either absolute or relative frequencies.

Author(s)

David Meyer <David.Meyer@R-project.org>

Examples

```
data("MSPatients")
independence_table(MSPatients)
independence_table(MSPatients, frequency = "relative")
```

JobSatisfaction *Job Satisfaction Data*

Description

Data from Petersen (1968) about the job satisfaction of 715 blue collar workers, selected from Danish Industry in 1968.

Usage

```
data("JobSatisfaction")
```

Format

A data frame with 8 observations and 4 variables.

Freq frequency.

management factor indicating quality of management (bad, good).

supervisor factor indicating supervisor's job satisfaction (low, high).

own factor indicating worker's own job satisfaction (low, high).

Source

E. B. Andersen (1991), The Statistical Analysis of Categorical Data, Table 5.4.

References

E. B. Andersen (1991), *The Statistical Analysis of Categorical Data*. 2nd edition. Springer-Verlag, Berlin.

E. Petersen (1968), *Job Satisfaction in Denmark*. (In Danish). Mentalhygiejnisk Forlag, Copenhagen.

Examples

```
data("JobSatisfaction")
structable(~ ., data = JobSatisfaction)

mantelhaen.test(xtabs(Freq ~ own + supervisor + management,
                      data = JobSatisfaction))
```

JointSports   *Opinions About Joint Sports*

Description

Data from a Danish study in 1983 and 1985 about sports activities and the opinion about joint sports with the other gender among 16–19 year old high school students.

Usage

```
data("JointSports")
```

Format

A data frame with 40 observations and 5 variables.

Freq frequency.

opinion factor indicating opinion about sports joint with the other gender (very good, good, indifferent, bad, very bad).

year factor indicating year of study (1983, 1985).

grade factor indicating school grade (1st, 3rd).

gender factor indicating gender (Boy, Girl).

Source

E. B. Andersen (1991), The Statistical Analysis of Categorical Data, page 210.

References

E. B. Andersen (1991), *The Statistical Analysis of Categorical Data*. 2nd edition. Springer-Verlag, Berlin.

Examples

```
library(MASS)
data("JointSports")
tab <- xtabs(Freq ~ gender + opinion + grade + year, data = JointSports)
doubledecker(opinion ~ gender + year + grade, data = tab)
loglm(~ opinion* (gender + grade+ year) + gender*year*grade, data = tab)
```

Kappa   *Cohen's Kappa and Weighted Kappa*

Description

Computes two agreement rates: Cohen's kappa and weighted kappa, and confidence bands.

Usage

```
Kappa(x, weights = c("Equal-Spacing", "Fleiss-Cohen"))
```

Arguments

x   a confusion matrix.

weights   either one of the character strings given in the default value, or a user-specified matrix with same dimensions as x.

Details

Cohen's kappa is the diagonal sum of the (possibly weighted) relative frequencies, corrected for expected values and standardized by its maximum value. The equal-spacing weights are defined by $1-|i-j|/(r-1)$, r number of columns/rows, and the Fleiss-Cohen weights by $1-|i-j|^2/(r-1)^2$. The latter one attaches greater importance to near disagreements.

Value

An object of class `"Kappa"` with three components:

Unweighted	numeric vector of length 2 with the kappa statistic (`value` component), along with Approximate Standard Error (`ASE` component)
Weighted	idem for the weighted kappa.
Weights	numeric matrix with weights used.

Note

The `summary` method also prints the weights.

There is a `confint` method for computing approximate confidence intervals.

Author(s)

David Meyer <David.Meyer@R-project.org>

References

Cohen, J. (1960), A coefficient of agreement for nominal scales. *Educational and Psychological Measurement*, **20**, 37–46.

Everitt, B.S. (1968), Moments of statistics kappa and weighted kappa. *The British Journal of Mathematical and Statistical Psychology*, **21**, 97–103.

Fleiss, J.L., Cohen, J., and Everitt, B.S. (1969), Large sample standard errors of kappa and weighted kappa. *Psychological Bulletin*, **72**, 332–327.

See Also

agreementplot, confint

Examples

```
data("SexualFun")
Kappa(SexualFun)
```

---

labeling_border   *Labeling Functions for Strucplots*

Description

These functions generate labeling functions used for strucplots.

Usage

```
labeling_border(labels = TRUE, varnames = labels,
  set_labels = NULL, set_varnames = NULL,
  tl_labels = NULL, alternate_labels = FALSE, tl_varnames = NULL,
  gp_labels = gpar(fontsize = 12),
  gp_varnames = gpar(fontsize = 12, fontface = 2),
  rot_labels = c(0, 90, 0, 90), rot_varnames = c(0, 90, 0, 90),
  pos_labels = "center", pos_varnames = "center",
  just_labels = "center", just_varnames = pos_varnames,
  boxes = FALSE, fill_boxes = FALSE,
  offset_labels = c(0, 0, 0, 0), offset_varnames = offset_labels,
  labbl_varnames = NULL, labels_varnames = FALSE, sep = ": ",
  abbreviate_labs = FALSE, rep = TRUE, clip = FALSE, ...)
labeling_values(value_type = c("observed", "expected", "residuals"),
                suppress = NULL, digits = 1, clip_cells = FALSE, ...)
labeling_residuals(suppress = NULL, digits = 1, clip_cells = FALSE, ...)
labeling_conditional(...)
labeling_left(rep = FALSE, pos_varnames = "left",
  pos_labels = "left", just_labels = "left", ...)
labeling_left2(tl_labels = TRUE, clip = TRUE, pos_varnames = "left",
  pos_labels = "left", just_labels = "left", ...)
labeling_cboxed(tl_labels = TRUE, boxes = TRUE, clip = TRUE,
  pos_labels = "center", ...)
labeling_lboxed(tl_labels = FALSE, boxes = TRUE, clip = TRUE,
  pos_labels = "left", just_labels = "left",
  labbl_varnames = FALSE, ...)
labeling_doubledecker(lab_pos = c("bottom", "top"), dep_varname = TRUE,
  boxes = NULL, clip = NULL, labbl_varnames = FALSE,
  rot_labels = rep.int(0, 4),
  pos_labels = c("left", "center", "left", "center"),
  just_labels = c("left", "left", "left", "center"),
  varnames = NULL, gp_varnames = gpar(fontsize = 12, fontface = 2),
  offset_varnames = c(0, -0.6, 0, 0), tl_labels = NULL, ...)
```

Arguments

labels   vector of logicals indicating whether labels should be drawn for a particular
                  dimension.

varnames	vector of logicals indicating whether variable names should be drawn for a particular dimension.
set_labels	An optional character vector with named components replacing the so-specified variable names. The component names must exactly match the variable names to be replaced.
set_varnames	An optional list with named components of character vectors replacing the labels of the so-specified variables. The component names must exactly match the variable names whose labels should be replaced.
tl_labels	vector of logicals indicating whether labels should be positioned on top (column labels) / left (row labels) for a particular dimension.
alternate_labels	
	vector of logicals indicating whether labels should be alternated on the top/bottom (left/right) side of the plot for a particular dimension.
tl_varnames	vector of logicals indicating whether variable names should be positioned on top (column labels) / on left (row labels) for a particular dimension.
gp_labels	list of objects of class "gpar" used for drawing the labels.
gp_varnames	list of objects of class "gpar" used for drawing the variable names.
rot_labels	vector of rotation angles for the labels for each of the four sides of the plot.
rot_varnames	vector of rotation angles for the variable names for each of the four sides of the plot.
pos_labels	character string of label positions ("left", "center", "right") for each of the variables.
pos_varnames	character string of variable names positions ("left", "center", "right") for each of the four sides of the plot.
just_labels	character string of label justifications ("left", "center", "right") for each of the variables.
just_varnames	character string of variable names justifications ("left", "center", "right") for each of the four sides of the plot.
boxes	vector of logicals indicating whether boxes should be drawn around the labels for a particular dimension.
fill_boxes	Either a vector of logicals, or a vector of characters, or a list of such vectors, specifying the fill colors for the boxes. "TRUE" and "FALSE" values are transformed into "grey" and "white", respectively. If fill_boxes is atomic, each component specifies a basic color for the corresponding dimension. This color is transformed into its HSV representation, and the value is varied from 50% to 100% to give a sequential color palette for the levels. For NA components, no palette is produced (no fill color). If fill_boxes is a list of vectors, each vector specifies the level colors of the corresponding dimension.
offset_labels, offset_varnames	
	numeric vector of length 4 indicating the offset of the labels (variable names) for each of the four sides of the plot.
labbl_varnames	vector of logicals indicating whether variable names should be drawn on the left (column variables) / on top (row variables) of the corresponding labels.

labels_varnames

        vector of logicals indicating, for each dimension, whether the variable name should be added to the corresponding labels or not.

sep   separator used if any component of `"labels_varnames"` is TRUE.

abbreviate_labs

        vector of integers or logicals indicating, for each dimension, the number of characters the labels should be abbreviated to. TRUE means 1 character, FALSE causes no abbreviation. Values are recycled as needed.

rep   vector of logicals indicating, for each dimension, whether labels should be repeated for all conditioning strata, or appear only once.

clip   vector of integers indicating, for each dimension, whether labels should be clipped to not overlap.

lab_pos   character string switching between `"top"` or `"bottom"` position of the labels (only used for `labeling_doubledecker`).

dep_varname   logical or character string. If logical, this is indicating whether the name of the dependent variable should be printed or not. A character string will be printed instead of the variable name taken from the dimnames.

value_type   character string specifying which values should be displayed in the cells.

suppress   numeric vector of length 2 specifying an interval of values that are not displayed. 0 values are never displayed. A single number, k, is treated as `c(-k, k)`. The default for labeling residuals is `c(-2,2)`. Use `suppress = 0` to show all non-zero values.

digits   integer specifying the number of digits used for rounding.

clip_cells   logical indicating whether the values should be clipped at the cell borders.

...   only used for `labeling_conditional` and `labeling_doubledecker`: parameters passed to `labeling_cells` and `labeling_border`.

Details

These functions generate labeling functions called by [strucplot](#) for their side-effect of adding labels to the plot. They suppose that a strucplot has been drawn and the corresponding viewport structure is pushed, since the positions of the viewports are used for the label positioning. Note that the functions can also be used 'stand-alone' as shown in the examples.

All values supplied to vectorized arguments can be 'abbreviated' by using named components which override the default component values. In addition, these defaults can be overloaded by the sequence of non-named components which are recycled as needed (see examples).

This help page only documents `labeling_border` and derived functions, more functions are described on the help page for [labeling_cells](#) and [labeling_list](#).

`labeling_left`, `labeling_left2`, `labeling_cboxed`, and `labeling_lboxed` are really just wrappers to `labeling_border`, and good examples for the parameter usage.

`labeling_residuals` is a trivial wrapper for `labeling_values`, which in turn calls `labeling_border` by additionally adding the observed or expected frequencies or residuals to the cells.

Value

A function with arguments:

d	`"dimnames"` attribute from the visualized contingency table, or the visualized table itself from which the `"dimnames"` attributes will then be extracted.
split_vertical	vector of logicals indicating the split directions.
condvars	integer vector of conditioning dimensions.

Author(s)

David Meyer <David.Meyer@R-project.org>

References

Meyer, D., Zeileis, A., and Hornik, K. (2006), The strucplot framework: Visualizing multi-way contingency tables with **vcd**. *Journal of Statistical Software*, **17(3)**, 1-48. URL http://www.jstatsoft.org/v17/i03/ and available as `vignette("strucplot")`.

See Also

[labeling_cells](#), [labeling_list](#), [structable](#), [grid.text](#)

Examples

```
data("Titanic")

mosaic(Titanic)

mosaic(Titanic, labeling = labeling_left)
labeling_left

mosaic(Titanic, labeling = labeling_cboxed)
labeling_cboxed

mosaic(Titanic, labeling = labeling_lboxed)
labeling_lboxed

data("PreSex")
mosaic(~ PremaritalSex + ExtramaritalSex | Gender + MaritalStatus,
  data = PreSex, labeling = labeling_conditional)

## specification of vectorized arguments
mosaic(Titanic, labeling_args = list(abbreviate_labs = c(Survived = TRUE)))

mosaic(Titanic, labeling_args = list(clip = TRUE, boxes = TRUE,
  fill_boxes = c(Survived = "green", "red")))

mosaic(Titanic, labeling_args = list(clip = TRUE, boxes = TRUE,
  fill_boxes = list(Sex = "red", "green")))

mosaic(Titanic, labeling_args = list(clip = TRUE, boxes = TRUE,
```

```
      fill_boxes = list(Sex = c(Male = "red", "blue"), "green")))

## change variable names
mosaic(Titanic, labeling_args = list(set_varnames = c(Sex = "Gender")))

## change labels
mosaic(Titanic, labeling_args = list(set_varnames = c(Survived = "Status"),
  set_labels = list(Survived = c("Survived", "Not Survived")), rep = FALSE))

## show frequencies
mosaic(Titanic, labeling = labeling_values)
```

---

labeling_cells_list   *Labeling Functions for Strucplots*

Description

These functions generate labeling functions that produce labels for strucplots.

Usage

```
labeling_cells(labels = TRUE, varnames = TRUE,
  abbreviate_labels = FALSE, abbreviate_varnames = FALSE,
  gp_text = gpar(), lsep = ": ", lcollapse = "\n",
  just = "center", pos = "center", rot = 0,
  margin = unit(0.5, "lines"), clip_cells = TRUE,
  text = NULL, ...)
labeling_list(gp_text = gpar(), just = "left", pos = "left", lsep = ": ",
  sep = " ", offset = unit(c(2, 2), "lines"),
  varnames = TRUE, cols = 2, ...)
```

Arguments

labels
:   vector of logicals indicating, for each dimension, whether labels for the factor levels should be drawn or not. Values are recycled as needed.

varnames
:   vector of logicals indicating, for each dimension, whether variable names should be drawn. Values are recycled as needed.

abbreviate_labels
:   vector of integers or logicals indicating, for each dimension, the number of characters the labels should be abbreviated to. TRUE means 1 character, FALSE causes no abbreviation. Values are recycled as needed.

abbreviate_varnames
:   vector of integers or logicals indicating, for each dimension, the number of characters the variable (i.e., dimension) names should be abbreviated to. TRUE means 1 character, FALSE causes no abbreviation. Values are recycled as needed.

gp_text
:   object of class "gpar" used for the text drawn.

lsep
:   character that separates variable names from the factor levels.

sep	character that separates the factor levels (only used for `labeling_list`).
offset	object of class `"unit"` of length 2 specifying the offset in x- and y-direction of the text block drawn under the strucplot (only used for `labeling_list`).
cols	number of text columns (only used for `labeling_list`).
lcollapse	character that separates several variable name/factor level-combinations. Typically a line break. (Only used for `labeling_cells`.)
just, pos	character string of length 1 (`labeling_list`) or at most 2 (`labeling_cells`) specifying the labels' horizontal position and justification (horizontal and vertical for `labeling_cells`).
rot	rotation angle in degrees, used for all labels (only used for `labeling_cells`).
margin	object of class `"unit"` (a numeric value is converted to `"lines"`) specifying an offset from the cell borders (only used for `labeling_cells`).
clip_cells	logical indicating whether text should be clipped at the cell borders (only used for `labeling_cells`).
text	Optionally, a character table of the same dimensions than the contingency table whose entries will then be used instead of the labels. `NA` entries are not drawn. This allows custom cell annotations (see examples). Only used for `labeling_cells`.
...	Currently not used.

Details

These functions generate labeling functions that can add different kinds of labels to an existing plot. Typically they are supplied to strucplot which then generates and calls the labeling function. They assume that a strucplot has been drawn and the corresponding viewport structure is pushed, so that by navigating through the viewport tree the labels can be positioned appropriately.

This help page only documents `labeling_list` and `labeling_cells`; more functions are described on the help page for labeling_border.

The functions can also be used 'stand-alone' as shown in the examples.

Using `labeling_list` will typically necessitate a bottom margin adjustment.

Value

A function with arguments:

d	`"dimnames"` attribute from the visualized contingency table, or the visualized table itself from which the `"dimnames"` attributes will then be extracted.
split_vertical	vector of logicals indicating the split directions.
condvars	integer vector of conditioning dimensions

Author(s)

David Meyer <David.Meyer@R-project.org>

References

Meyer, D., Zeileis, A., and Hornik, K. (2006), The strucplot framework: Visualizing multi-way contingency tables with **vcd**. *Journal of Statistical Software*, **17(3)**, 1-48. URL http://www.jstatsoft.org/v17/i03/ and available as vignette(″strucplot″).

See Also

[labeling_border](), [structable](), [grid.text]()

Examples

```
data("Titanic")

mosaic(Titanic, labeling = labeling_cells)
mosaic(Titanic, labeling = labeling_list)

## A more complex example, adding the observed frequencies
## to a mosaic plot:
tab <- ifelse(Titanic < 6, NA, Titanic)
mosaic(Titanic, pop = FALSE)
labeling_cells(text = tab, margin = 0)(Titanic)
```

---

legends   *Legend Functions for Strucplots*

Description

These functions generate legend functions for residual-based shadings.

Usage

```
legend_resbased(fontsize = 12, fontfamily = ″″,
  x = unit(1, ″lines″), y = unit(0.1,″npc″),
  height = unit(0.8, ″npc″),
  width = unit(0.7, ″lines″),
  digits = 2, check_overlap = TRUE, text = NULL,
  steps = 200, ticks = 10, pvalue = TRUE,  range = NULL)
legend_fixed(fontsize = 12, fontfamily = ″″, x = unit(1, ″lines″), y = NULL,
  height = NULL, width = unit(1.5, ″lines″), steps = 200,
  digits = 1, space = 0.05, text = NULL, range = NULL)
```

Arguments

fontsize	fontsize of title and p-value text.
fontfamily	fontfamily of all text.
x, y	objects of class ″unit″ indicating the coordinates of the title. For legend_fixed, the default for y is computed as to leave enough space for the specified text.

height, width	object of class "unit" indicating the height/width of the legend. For legend_fixed, the default for y is computed as to align upper margins of legend and actual plot.
digits	number of digits for the scale labels.
check_overlap	logical indicating whether overlap of scale labels should be inhibited or not.
space	For legend_fixed only: proportion of space between the tiles.
text	character string indicating the title of the legend.
steps	granularity of the color gradient.
ticks	number of scale ticks.
pvalue	logical indicating whether the p-value should be visualized under the scale or not.
range	Numeric vector of length 2 for setting the legend range. Computed from the residuals if omitted. NA values are replaced by the corresponding minimum / maximum of the residuals.

Details

These functions generate legend functions for residual-based shadings, visualizing deviations from expected values of an hypothesized independence model. Therefore, the legend uses a supplied shading function to visualize the color gradient for the residuals range. legend_fixed is inspired by the legend used in mosaicplot. For more details on the shading functions and their return values, see shadings.

Value

A function with arguments:

residuals	residuals from the fitted independence model to be visualized.
shading	shading function computing colors from residuals (see details).
autotext	character vector indicating the title to be used when no text argument is specified. Allows strucplot to generate sensible defaults depending on the residuals type.

Author(s)

David Meyer <David.Meyer@R-project.org>

References

Meyer, D., Zeileis, A., and Hornik, K. (2006), The strucplot framework: Visualizing multi-way contingency tables with **vcd**. *Journal of Statistical Software*, **17(3)**, 1-48. URL http://www.jstatsoft.org/v17/i03/ and available as vignette("strucplot").

Meyer, D., Zeileis, A., Hornik, K. (2003), Visualizing independence using extended association plots. *Proceedings of the 3rd International Workshop on Distributed Statistical Computing*, K. Hornik, F. Leisch, A. Zeileis (eds.), ISSN 1609-395X. http://www.ci.tuwien.ac.at/Conferences/DSC-2003/Proceedings/

See Also

structable, shadings

Examples

```
data("Titanic")

mosaic(Titanic, shade = TRUE, legend = legend_resbased)
mosaic(Titanic, shade = TRUE, legend = legend_fixed, gp = shading_Friendly)
```

Lifeboats   *Lifeboats on the Titanic*

Description

Data from Mersey (1912) about the 18 (out of 20) lifeboats launched before the sinking of the S. S. Titanic.

Usage

```
data("Lifeboats")
```

Format

A data frame with 18 observations and 8 variables.

launch launch time in "POSIXt" format.

side factor. Side of the boat.

boat factor indicating the boat.

crew number of male crew members on board.

men number of men on board.

women number of women (including female crew) on board.

total total number of passengers.

cap capacity of the boat.

Source

M. Friendly (2000), Visualizing Categorical Data: http://euclid.psych.yorku.ca/ftp/sas/vcd/catdata/lifeboat.sas

References

L. Mersey (1912), Report on the loss of the "Titanic" (S. S.). Parliamentary command paper 6452.

M. Friendly (2000), *Visualizing Categorical Data*. SAS Institute, Cary, NC.

Examples

```
data("Lifeboats")
attach(Lifeboats)
ternaryplot(
  Lifeboats[,4:6],
  pch = ifelse(side == "Port", 1, 19),
  col = ifelse(side == "Port", "red", "blue"),
  id  = ifelse(men / total > 0.1, as.character(boat), NA),
  prop_size = 2,
  dimnames_position = "edge",
  main = "Lifeboats on the Titanic"
)
grid_legend(0.8, 0.9, c(1, 19), c("red", "blue"),
  c("Port", "Starboard"), title = "SIDE")
detach(Lifeboats)
```

---

mar_table   *Table with Marginal Sums*

Description

Adds row and column sums to a two-way table.

Usage

```
mar_table(x)
```

Arguments

x   a two-way table.

Value

A table with row and column totals added.

Author(s)

David Meyer <David.Meyer@R-project.org>

Examples

```
data("SexualFun")
mar_table(SexualFun)
```

---

mosaic   *Extended Mosaic Plots*

Description

Plots (extended) mosaic displays.

Usage

```
## Default S3 method:
mosaic(x, condvars = NULL,
  split_vertical = NULL, direction = NULL, spacing = NULL,
  spacing_args = list(), gp = NULL, expected = NULL, shade = NULL,
 highlighting = NULL, highlighting_fill = grey.colors, highlighting_direction = NULL,
  zero_size = 0.5, zero_split = FALSE, zero_shade = NULL,
  zero_gp = gpar(col = 0), panel = NULL, main = NULL, sub = NULL, ...)
## S3 method for class 'formula'
mosaic(formula, data, highlighting = NULL,
  ..., main = NULL, sub = NULL, subset = NULL, na.action = NULL)
```

Arguments

x	a contingency table in array form, with optional category labels specified in the `dimnames(x)` attribute, or an object of class `"structable"`.
condvars	vector of integers or character strings indicating conditioning variables, if any. The table will be permuted to order them first.
formula	a formula specifying the variables used to create a contingency table from `data`. For convenience, conditioning formulas can be specified; the conditioning variables will then be used first for splitting. If any, a specified response variable will be highlighted in the cells.
data	either a data frame, or an object of class `"table"` or `"ftable"`.
subset	an optional vector specifying a subset of observations to be used.
na.action	a function which indicates what should happen when the data contain NAs. Ignored if `data` is a contingency table.
zero_size	size of the bullets used for zero entries (if 0, no bullets are drawn).
zero_split	logical controlling whether zero cells should be further split. If FALSE and `zero_shade` is FALSE, only one bullet is drawn (centered) for unsplit zero cells. If FALSE and `zero_shade` is TRUE, a bullet for each zero cell is drawn to allow, e.g., residual-based shadings to be effective also for zero cells.
zero_shade	logical controlling whether zero bullets should be shaded. The default is TRUE if `shade` is TRUE or `expected` is not null or `gp` is not null, and FALSE otherwise.
zero_gp	object of class `"gpar"` used for zero bullets in case they are *not* shaded.

split_vertical	vector of logicals of length k, where k is the number of margins of x (default: FALSE). Values are recycled as needed. A TRUE component indicates that the tile(s) of the corresponding dimension should be split vertically, FALSE means horizontal splits. Ignored if direction is not NULL.
direction	character vector of length k, where k is the number of margins of x (values are recycled as needed). For each component, a value of "h" indicates that the tile(s) of the corresponding dimension should be split horizontally, whereas "v" indicates vertical split(s).
spacing	spacing object, spacing function, or corresponding generating function (see strucplot for more information). The default is spacing_equal if x has two dimensions, spacing_increase for more dimensions, and spacing_conditional if conditioning variables are specified using condvars or the formula interface.
spacing_args	list of arguments for the generating function, if specified (see strucplot for more information).
gp	object of class "gpar", shading function or a corresponding generating function (see details and shadings). Components of "gpar" objects are recycled as needed along the last splitting dimension. Ignored if shade = FALSE.
shade	logical specifying whether gp should be used or not (see gp). If TRUE and expected is unspecified, a default model is fitted: if condvars (see strucplot) is specified, a corresponding conditional independence model, and else the total independence model.
expected	optionally, an array of expected values of the same dimension as x, or alternatively the corresponding independence model specification as used by loglin or loglm (see strucplot).
highlighting	character vector or integer specifying a variable to be highlighted in the cells.
highlighting_fill	
	color vector or palette function used for a highlighted variable, if any.
highlighting_direction	
	Either "left", "right", "top", or "bottom" specifying the direction of highlighting in the cells.
panel	Optional function with arguments: residuals, observed, expected, index, gp, and name called by the struc_mosaic workhorse for each tile that is drawn in the mosaic. index is an integer vector with the tile's coordinates in the contingency table, gp a gpar object for the tile, and name a label to be assigned to the drawn grid object.
main, sub	either a logical, or a character string used for plotting the main (sub) title. If logical and TRUE, the name of the data object is used.
...	Other arguments passed to strucplot

Details

Mosaic displays have been suggested in the statistical literature by Hartigan and Kleiner (1984) and have been extended by Friendly (1994). mosaicplot is a base graphics implementation and mosaic is a much more flexible and extensible grid implementation.

`mosaic` is a generic function which currently has a default method and a formula interface. Both are high-level interfaces to the `strucplot` function, and produce (extended) mosaic displays. Most of the functionality is described there, such as specification of the independence model, labeling, legend, spacing, shading, and other graphical parameters.

A mosaic plot is an area proportional visualization of a (possibly higher-dimensional) table of expected frequencies. It is composed of tiles (corresponding to the cells) created by recursive vertical and horizontal splits of a square. The area of each tile is proportional to the corresponding cell entry, *given* the dimensions of previous splits.

An *extended* mosaic plot, in addition, visualizes the fit of a particular log-linear model. Typically, this is done by residual-based shadings where color and/or outline of the tiles visualize sign, size and possibly significance of the corresponding residual.

The layout is very flexible: the specification of shading, labeling, spacing, and legend is modularized (see `strucplot` for details).

In contrast to the `mosaicplot` function in **graphics**, the splits start with the *horizontal* direction by default to match the printed output of `structable`.

Value

The `"structable"` visualized is returned invisibly.

Author(s)

David Meyer <David.Meyer@R-project.org>

References

Hartigan, J.A., and Kleiner, B. (1984), A mosaic of television ratings. *The American Statistician*, **38**, 32–35.

Emerson, J. W. (1998), Mosaic displays in S-PLUS: A general implementation and a case study. *Statistical Computing and Graphics Newsletter (ASA)*, **9**, 1, 17–23.

Friendly, M. (1994), Mosaic displays for multi-way contingency tables. *Journal of the American Statistical Association*, **89**, 190–200.

Meyer, D., Zeileis, A., and Hornik, K. (2006), The strucplot framework: Visualizing multi-way contingency tables with **vcd**. *Journal of Statistical Software*, **17(3)**, 1-48. URL http://www.jstatsoft.org/v17/i03/ and available as `vignette("strucplot", package = "vcd")`.

The home page of Michael Friendly (http://datavis.ca) provides information on various aspects of graphical methods for analyzing categorical data, including mosaic plots. In particular, there are many materials for his course "Visualizing Categorical Data with SAS and R" at http://datavis.ca/courses/VCD/.

See Also

`assoc`, `strucplot`, `mosaicplot`, `structable`, `doubledecker`

Examples

```
library(MASS)
data("Titanic")
mosaic(Titanic)

## Formula interface for tabulated data plus shading and legend:
mosaic(~ Sex + Age + Survived, data = Titanic,
  main = "Survival on the Titanic", shade = TRUE, legend = TRUE)

data("HairEyeColor")
mosaic(HairEyeColor, shade = TRUE)
## Independence model of hair and eye color and sex.  Indicates that
## there are significantly more blue eyed blond females than expected
## in the case of independence (and too few brown eyed blond females).

mosaic(HairEyeColor, shade = TRUE, expected = list(c(1,2), 3))
## Model of joint independence of sex from hair and eye color.  Males
## are underrepresented among people with brown hair and eyes, and are
## overrepresented among people with brown hair and blue eyes, but not
## "significantly".

## Formula interface for raw data: visualize crosstabulation of numbers
## of gears and carburettors in Motor Trend car data.
data("mtcars")
mosaic(~ gear + carb, data = mtcars, shade = TRUE)

data("PreSex")
mosaic(PreSex, condvars = c(1,4))
mosaic(~ ExtramaritalSex + PremaritalSex | MaritalStatus + Gender,
       data = PreSex)

## Highlighting:
mosaic(Survived ~ ., data = Titanic)

data("Arthritis")
mosaic(Improved ~ Treatment | Sex, data = Arthritis, zero_size = 0)
mosaic(Improved ~ Treatment | Sex, data = Arthritis, zero_size = 0,
       highlighting_direction = "right")
```

MSPatients   *Diagnosis of Multiple Sclerosis*

Description

Data from Westlund \& Kurland (1953) on the diagnosis of multiple sclerosis (MS): two samples of patients, one from Winnipeg and one from New Orleans, were each rated by two neurologists (one from each city) in four diagnostic categories.

Usage

```
data("MSPatients")
```

Format

A 3-dimensional array resulting from cross-tabulating 218 observations on 3 variables. The variables and their levels are as follows:

No	Name	Levels
1	New Orleans Neurologist	Certain, Probable, Possible, Doubtful
2	Winnipeg Neurologist	Certain, Probable, Possible, Doubtful
3	Patients	Winnipeg, New Orleans

Source

M. Friendly (2000), Visualizing Categorical Data: http://euclid.psych.yorku.ca/ftp/sas/vcd/catdata/msdiag.sas

References

K. B. Westlund \& L. T. Kurland (1953), Studies on multiple sclerosis in Winnipeg, Manitoba and New Orleans, Louisiana, *American Journal of Hygiene*, **57**, 380–396.

M. Friendly (2000), *Visualizing Categorical Data*. SAS Institute, Cary, NC.

Examples

```
data("MSPatients")
## best visualized using a resized device, e.g. using:
## get(getOption("device"))(width = 12)
pushViewport(viewport(layout = grid.layout(ncol = 2)))
pushViewport(viewport(layout.pos.col = 1))
agreementplot(t(MSPatients[,,1]), main = "Winnipeg Patients",
              newpage = FALSE)
popViewport()
pushViewport(viewport(layout.pos.col = 2))
agreementplot(t(MSPatients[,,2]), main = "New Orleans Patients",
              newpage = FALSE)
popViewport(2)
dev.off()
```

NonResponse	*Non-Response Survey Data*

Description

Data about non-response for a Danish survey in 1965.

Usage

```
data("NonResponse")
```

Format

A data frame with 12 observations and 4 variables.

Freq frequency.

residence factor indicating whether residence was in Copenhagen, in a city outside Copenhagen or at the countryside (Copenhagen, City, Country).

response factor indicating whether a response was given (yes, no).

gender factor indicating gender (male, female).

Source

E. B. Andersen (1991), The Statistical Analysis of Categorical Data, Table 5.17.

References

E. B. Andersen (1991), *The Statistical Analysis of Categorical Data*. 2nd edition. Springer-Verlag, Berlin.

Examples

```
data("NonResponse")
structable(~ ., data = NonResponse)
```

---

oddsratio   *(Log) Odds Ratios*

Description

Computes (log) odds ratios and their asymptotic standard errors for (possibly) stratified data.

Usage

```
oddsratio(x, stratum = NULL, log = TRUE)
## S3 method for class 'oddsratio'
plot(x, conf_level = 0.95, type = "o",
  xlab = NULL, ylab = NULL, xlim = NULL, ylim = NULL, whiskers = 0.1,
  baseline = TRUE, transpose = FALSE, ...)
```

Arguments

x	a 2 by 2 by . . . table.
stratum	vector of strata dimensions.
log	if FALSE, ordinary odds ratios are computed.
conf_level	if not NULL or FALSE, conf_level-% confidence intervals are plotted for each data point.
type	plot type.
xlab	label for the x-axis. Defaults to "Strata" if transpose is FALSE.
ylab	label for the y-axis. Defaults to "Strata" if transpose is TRUE.
xlim	x-axis limits. Ignored if transpose is FALSE.
ylim	y-axis limits. Ignored if transpose is TRUE.
baseline	if TRUE, a red dashed line is plotted at a value of 1 (in case of odds) or 0 (in case of log-odds).
transpose	if TRUE, the plot is transposed.
whiskers	width of the confidence interval whiskers.
...	other graphics parameters (see [par](#)).

Value

An object of class "logoddsratio", which is simply a vector of (log) odds ratios with dimensionality depending on stratum, along with the following attributes:

ASE	a numeric vector with the asymptotic standard errors.
log	logical indicating whether log odds ratios or common odds ratios are computed.

Note

In case of zero entries, 0.5 will be added to the table.

The summary method prints the standard errors and—for log odds ratios—also computes and prints asymptotic z tests (standardized log odds ratios) and the corresponding p values.

There is a confint method for computing confidence intervals for the (log) odds ratios.

The plot method plots (log) odds ratios, computed by oddsratio for $2 \times 2 \times k$ tables, along with confidence intervals.

Author(s)

David Meyer <David.Meyer@R-project.org>

References

M. Friendly (2000), *Visualizing Categorical Data*. SAS Institute, Cary, NC.

See Also

[confint](#)

Examples

```
## load Coal Miners data
data("CoalMiners")

## compute log odds ratios
lor <- oddsratio(CoalMiners)
lor

## summary with z tests
summary(lor)

## confidence intervals
confint(lor)

## visualization
plot(lor,
     xlab = "Age Group",
     main = "Breathelessness and Wheeze in Coal Miners")

## add linear model
age <- seq(20, 60, by = 5)
m <- lm(lor ~ age)
lines(fitted(m), col = "red", lwd=2)
```

Ord_plot   *Ord Plots*

Description

Ord plots for diagnosing discrete distributions.

Usage

```
Ord_plot(obj, legend = TRUE, estimate = TRUE, tol = 0.1, type = NULL,
  xlim = NULL, ylim = NULL, xlab = "Number of occurrences",
  ylab = "Frequency ratio", main = "Ord plot", gp = gpar(cex = 0.5),
  lwd = c(2,2), lty=c(2,1), col=c("black", "red"),
  name = "Ord_plot", newpage = TRUE, pop = TRUE, ...)
Ord_estimate(x, type = NULL, tol = 0.1)
```

Arguments

obj	either a vector of counts, a 1-way table of frequencies of counts or a data frame or matrix with frequencies in the first column and the corresponding counts in the second column.
legend	logical. Should a legend be plotted?.
estimate	logical. Should the distribution and its parameters be estimated from the data? See details.

tol	tolerance for estimating the distribution. See details.
type	a character string indicating the distribution, must be one of `"poisson"`, `"binomial"`, `"nbinomial"` or `"log-series"` or NULL. In the latter case the distribution is estimated from the data. See details.
xlim	limits for the x axis.
ylim	limits for the y axis.
xlab	a label for the x axis.
ylab	a label for the y axis.
main	a title for the plot.
gp	a `"gpar"` object controlling the grid graphical parameters of the points.
lwd, lty	vectors of length 2, giving the line width and line type used for drawing the OLS line and the WLS lines.
col	vector of length 2 giving the colors used for drawing the OLS and WLS lines.
name	name of the plotting viewport.
newpage	logical. Should `grid.newpage` be called before plotting?
pop	logical. Should the viewport created be popped?
...	further arguments passed to `grid.points`.
x	a vector giving intercept and slope for the (fitted) line in the Ord plot.

Details

The Ord plot plots the number of occurrences against a certain frequency ratio (see Friendly (2000) for details) and should give a straight line if the data comes from a poisson, binomial, negative binomial or log-series distribution. The intercept and slope of this straight line conveys information about the underlying distribution.

`Ord_plot` fits a usual OLS line (black) and a weighted OLS line (red). From the coefficients of the latter the distribution is estimated by `Ord_estimate` as described in Table 2.10 in Friendly (2000). To judge whether a coefficient is positive or negative a tolerance given by `tol` is used. If none of the distributions fits well, no parameters are estimated. Be careful with the conclusions from `Ord_estimate` as it implements just some simple heuristics!

Value

A vector giving the intercept and slope of the weighted OLS line.

Author(s)

Achim Zeileis <Achim.Zeileis@R-project.org>

References

J. K. Ord (1967), Graphical methods for a class of discrete distributions, *Journal of the Royal Statistical Society*, **A 130**, 232–238.

Michael Friendly (2000), *Visualizing Categorical Data*. SAS Institute, Cary, NC.

Examples

```
## Simulated data examples:
dummy <- rnbinom(1000, size = 1.5, prob = 0.8)
Ord_plot(dummy)

## Real data examples:
data("HorseKicks")
data("Federalist")
data("Butterfly")
data("WomenQueue")

grid.newpage()
pushViewport(viewport(layout = grid.layout(2, 2)))

pushViewport(viewport(layout.pos.col=1, layout.pos.row=1))
Ord_plot(HorseKicks, main = "Death by horse kicks", newpage = FALSE)
popViewport()

pushViewport(viewport(layout.pos.col=1, layout.pos.row=2))
Ord_plot(Federalist, main = "Instances of 'may' in Federalist papers", newpage = FALSE)
popViewport()

pushViewport(viewport(layout.pos.col=2, layout.pos.row=1))
Ord_plot(Butterfly, main = "Butterfly species collected in Malaya", newpage = FALSE)
popViewport()

pushViewport(viewport(layout.pos.col=2, layout.pos.row=2))
Ord_plot(WomenQueue, main = "Women in queues of length 10", newpage = FALSE)
popViewport(2)
```

OvaryCancer   *Ovary Cancer Data*

Description

Data from Obel (1975) about a retrospective study of ovary cancer carried out in 1973. Information was obtained from 299 women, who were operated for ovary cancer 10 years before.

Usage

```
data("OvaryCancer")
```

Format

A data frame with 16 observations and 5 variables.

Freq frequency.

stage factor indicating the stage of the cancer at the time of operation (early, advanced).

operation factor indicating type of operation (radical, limited).

survival factor indicating survival status after 10 years (yes, no).

xray factor indicating whether X-ray treatment was received (yes, no).

Source

E. B. Andersen (1991), The Statistical Analysis of Categorical Data, Table 6.4.

References

E. B. Obel (1975), A Comparative Study of Patients with Cancer of the Ovary Who Have Survived More or Less Than 10 Years. *Acta Obstetricia et Gynecologica Scandinavica*, **55**, 429-439.

E. B. Andersen (1991), *The Statistical Analysis of Categorical Data*. 2nd edition. Springer-Verlag, Berlin.

Examples

```
data("OvaryCancer")
tab <- xtabs(Freq ~ xray + survival + stage + operation, data = OvaryCancer)
ftable(tab, col.vars = "survival", row.vars = c("stage", "operation", "xray"))

## model: ~ xray * operation * stage + survival * stage
## interpretation: treat xray, operation, stage as fixed margins,
##    the survival depends on stage, but not xray and operation.
doubledecker(survival ~ stage + operation + xray, data = tab)
mosaic(~ stage + operation + xray + survival,
  split = c(FALSE, TRUE, TRUE, FALSE), data = tab, keep = FALSE,
  gp = gpar(fill = rev(grey.colors(2))))
mosaic(~ stage + operation + xray + survival,
  split = c(FALSE, TRUE, TRUE, FALSE), data = tab, keep = FALSE,
  expected = ~ xray * operation * stage + survival*stage)
```

Pairs plot panel functions for diagonal cells
                Diagonal Panel Functions for Table Pairs Plot

---

Description

Diagonal panel functions for pairs.table.

Usage

```
pairs_barplot(gp_bars = NULL,
  gp_vartext = gpar(fontsize = 17),
  gp_leveltext = gpar(),
  just_leveltext = c("center", "bottom"),
  just_vartext = c("center", "top"),
  rot = 0, abbreviate = FALSE, check_overlap = TRUE, fill = "grey",
  var_offset = unit(1, "npc"), ...)
```

```
pairs_text(dimnames = TRUE, gp_vartext = gpar(fontsize = 17),
  gp_leveltext = gpar(), gp_border = gpar(), ...)
pairs_diagonal_text(varnames = TRUE, gp_vartext = gpar(fontsize = 17, fontface = "bold"),
  gp_leveltext = gpar(), gp_border = gpar(), pos = c("right","top"),
  distribute = c("equal","margin"), rot = 0, ...)
pairs_diagonal_mosaic(split_vertical = TRUE, margins = unit(0, "lines"),
  offset_labels = -0.4, offset_varnames = 0,
  gp = NULL, fill = "grey", labeling = labeling_values, alternate_labels = TRUE, ...)
```

Arguments

dimnames	vector of logicals indicating whether the factor levels should be displayed (only used for `pairs_text`).
varnames	vector of logicals indicating whether the variable names should be displayed (only used for `pairs_text_diagonal`).
gp_bars	object of class "gpar" used for bars (only used for `pairs_barplot`). If unspecified, the default is to set the `fill` component of this object to the `fill` argument.
gp_vartext	object of class "gpar" used for the factor names.
gp_leveltext	object of class "gpar" used for the factor levels.
gp_border	object of class "gpar" used for the border (only used for `pairs_text`).
gp	object of class "gpar" used for the tiles (only used for `pairs_diagonal_mosaic`). If unspecified, the default is to set the `fill` component of this object to the `fill` argument.
fill	color vector or palette function used for the fill colors of bars (for `pairs_barplot`) or tiles (for `pairs_diagonal_mosaic`).
labeling	labeling function, passed to `mosaic()`
alternate_labels	should labels alternate top/bottom?
just_leveltext, just_vartext	character string indicating the justification for variable names and levels.
pos	character string of length 2 controlling the horizontal and vertical position of the variable names (only used for `pairs_text_diagonal`).
rot	rotation angle for the variable levels.
distribute	character string indicating whether levels should be distributed equally or according to the margins (only used for `pairs_text_diagonal`).
abbreviate	integer or logical indicating the number of characters the labels should be abbreviated to. TRUE means 1 character, FALSE causes no abbreviation.
check_overlap	If TRUE, some levels will suppressed to avoid overlapping, if any.
split_vertical	vector of logicals of length k, where k is the number of margins of x (values are recycled as needed). A TRUE component indicates that the tile(s) of the corresponding dimension should be split vertically, FALSE means horizontal splits. Default is FALSE.

margins   either an object of class `"unit"` of length 4, or a numeric vector of length 4.
                    The elements are recycled as needed. The four components specify the top,
                    right, bottom, and left margin of the plot, respectively. When a numeric vec-
                    tor is supplied, the numbers are interpreted as `"lines"` units. In addition, the
                    unit or numeric vector may have named arguments ('`top`', '`right`', '`bottom`',
                    and '`left`'), in which case the non-named arguments specify the default values
                    (recycled as needed), overloaded by the named arguments.

offset_labels, offset_varnames
                    numeric vector of length 4 indicating the offset of the labels (variable names)
                    for each of the four sides of the plot.

var_offset   object of class `"unit"` specifying the offset of variable names from the bottom
                    of the bar plots created by pairs_barplot. If numeric, the unit defaults to
                    `"npc"`.

...   other parameters passed to the underlying graphics functions.

Details

In the diagonal cells, the pairsplot visualizes statistics or information for each dimension (that is: the
single factors) alone. `pairs_text` displays the factor's name, and optionally also the factor levels.
`pairs_barplot` produces a bar plot of the corresponding factor, along with the factor's name.

Value

A function with one argument: the marginal table for the corresponding dimension.

Author(s)

David Meyer <David.Meyer@R-project.org>

See Also

`pairs.table`, `pairs_assoc`, `pairs_mosaic`

Examples

```
data("UCBAdmissions")

pairs(UCBAdmissions) # pairs_barplot is default
pairs(UCBAdmissions, diag_panel = pairs_text)

pairs(UCBAdmissions, diag_panel = pairs_diagonal_text)
pairs(Titanic, diag_panel = pairs_diagonal_text)
pairs(Titanic, diag_panel = pairs_diagonal_text(distribute = "margin"))
pairs(Titanic,
      diag_panel = pairs_diagonal_text(distribute = "margin", rot = 45))
```

Pairs plot panel functions for off-diagonal cells
Off-diagonal Panel Functions for Table Pairs Plot

Description

Off-diagonal panel functions for `pairs.table`.

Usage

```
pairs_strucplot(panel = mosaic,
  type = c("pairwise", "total", "conditional", "joint"),
  legend = FALSE, margins = c(0, 0, 0, 0), labeling = NULL, ...)
pairs_assoc(...)
pairs_mosaic(...)
pairs_sieve(...)
```

Arguments

panel	function to be used for the plots in each cell, such as `pairs_assoc`, `pairs_mosaic`, and `pairs_sieve`.
type	character string specifying the type of independence model visualized in the cells.
legend	logical specifying whether a legend should be displayed in the cells or not.
margins	margins inside each cell (see `strucplot`).
labeling	labeling function or labeling-generating function (see `strucplot`).
...	`pairs_mosaic`, `pairs_assoc`, and `pairs_sieve`: parameters passed to `pairs_strucplot`. `pairs_strucplot`: other parameters passed to panel function.

Details

These functions really just wrap `assoc`, `sieve`, and `mosaic` by basically inhibiting labeling and legend-drawing and setting the margins to 0.

Value

A function with arguments:

x	contingency table.
i, j	cell coordinates.

Author(s)

David Meyer <David.Meyer@R-project.org>

References

Cohen, A. (1980), On the graphical display of the significant components in a two-way contingency table. *Communications in Statistics—Theory and Methods*, **A9**, 1025–1041.

Friendly, M. (1992), Graphical methods for categorical data. *SAS User Group International Conference Proceedings*, **17**, 190–200. <http://datavis.ca/sugi/sugi17.pdf>

See Also

`pairs.table`, `pairs_text`, `pairs_barplot`, `assoc`, `mosaic`

Examples

```
data("UCBAdmissions")
data("PreSex")

pairs(PreSex)
pairs(UCBAdmissions)
pairs(UCBAdmissions, upper_panel_args = list(shade = FALSE))
pairs(UCBAdmissions, lower_panel = pairs_mosaic(type = "conditional"))
pairs(UCBAdmissions, upper_panel = pairs_assoc)
```

---

pairs.table	*Pairs Plot for Contingency Tables*

Description

Produces a matrix of strucplot displays.

Usage

```
## S3 method for class 'table'
pairs(x, upper_panel = pairs_mosaic, upper_panel_args = list(),
  lower_panel = pairs_mosaic, lower_panel_args = list(),
  diag_panel = pairs_diagonal_mosaic, diag_panel_args = list(),
  main = NULL, sub = NULL, main_gp = gpar(fontsize = 20),
  sub_gp = gpar(fontsize = 15), space = 0.3,
  newpage = TRUE, pop = TRUE, margins = unit(1, "lines"), ...)
```

Arguments

x	a contingency table in array form, with optional category labels specified in the `dimnames(x)` attribute.
upper_panel	function for the upper triangle of the matrix, or corresponding generating function. If `NULL`, no panel is drawn.
upper_panel_args	
	list of arguments for the generating function, if specified.

lower_panel	function for the lower triangle of the matrix, or corresponding generating function. If NULL, no panel is drawn.
lower_panel_args	
	list of arguments for the panel-generating function, if specified.
diag_panel	function for the diagonal of the matrix, or corresponding generating function. If NULL, no panel is drawn.
diag_panel_args	
	list of arguments for the generating function, if specified.
main	either a logical, or a character string used for plotting the main title. If main is a logical and TRUE, the name of the object supplied as x is used.
sub	a character string used for plotting the subtitle. If sub is a logical and TRUE and main is unspecified, the name of the object supplied as x is used.
main_gp, sub_gp	
	object of class "gpar" containing the graphical parameters used for the main (sub) title, if specified.
space	double specifying the distance between the cells.
newpage	logical controlling whether a new grid page should be created.
pop	logical indicating whether all viewports should be popped after the plot has been drawn.
margins	either an object of class "unit" of length 4, or a numeric vector of length 4. The elements are recycled as needed. The four components specify the top, right, bottom, and left margin of the plot, respectively. When a numeric vector is supplied, the numbers are interpreted as "lines" units. In addition, the unit or numeric vector may have named arguments ('top', 'right', 'bottom', and 'left'), in which case the non-named arguments specify the default values (recycled as needed), overloaded by the named arguments.
...	For convenience, list of arguments for the panel-generating functions of upper and lower panels, if specified.

Details

This is a [pairs](#) method for objects inheriting from class "table" or "structable". It plots a matrix of pairwise mosaic plots.

Four independence types are distinguished: "pairwise", "total", "conditional" and "joint". The pairwise mosaic matrix shows bivariate marginal relations, collapsed over all other variables. The total independence mosaic matrix shows mosaic plots for mutual independence, i.e., for marginal and conditional independence among all pairs of variables. The conditional independence mosaic matrix shows mosaic plots for conditional independence for each pair of variables, given all other variables. The joint independence mosaic matrix shows mosaic plots for joint independence of all pairs of variables from the others.

This method uses panel functions called for each cell of the matrix which can be different for upper matrix, lower matrix, and diagonal cells. Correspondingly, for each panel parameter *foo* (= 'upper', 'lower', or 'diag'), pairs.table takes two arguments: *foo_panel* and *foo_panel_args*, which can be used to specify the parameters as follows:

1. Passing a suitable panel function to *foo_panel* which subsequently is called for each cell with the corresponding coordinates.

2. Passing a corresponding *generating function* (of class "panel_generator") to *foo_panel*, along with parameters passed to *foo_panel_args*, that generates such a function.

Hence, the second approach is equivalent to the first if *foo_panel(foo_panel_args)* is passed to *foo_panel*.

Author(s)

David Meyer <David.Meyer@R-project.org>

References

Cohen, A. (1980), On the graphical display of the significant components in a two-way contingency table. *Communications in Statistics—Theory and Methods*, **A9**, 1025–1041.

Friendly, M. (1992), Graphical methods for categorical data. *SAS User Group International Conference Proceedings*, **17**, 190–200. http://datavis.ca/sugi/sugi17.pdf

Meyer, D., Zeileis, A., and Hornik, K. (2006), The strucplot framework: Visualizing multi-way contingency tables with **vcd**. *Journal of Statistical Software*, **17(3)**, 1-48. URL http://www.jstatsoft.org/v17/i03/ and available as vignette("strucplot").

See Also

pairs_mosaic, pairs_assoc, pairs_sieve, pairs_diagonal_text, pairs_diagonal_mosaic, pairs_text, pairs_barplot, assoc, sieve, mosaic

Examples

```
data("UCBAdmissions")
data("PreSex")
data(HairEyeColor)
hec = structable(Eye ~ Sex + Hair, data = HairEyeColor)

pairs(PreSex)
pairs(UCBAdmissions)
pairs(UCBAdmissions, upper_panel_args = list(shade = TRUE))
pairs(UCBAdmissions, lower_panel = pairs_mosaic(type = "conditional"))
pairs(UCBAdmissions, diag_panel = pairs_text)
pairs(UCBAdmissions, upper_panel = pairs_assoc, shade = TRUE)
pairs(hec, highlighting = 2, diag_panel_args = list(fill = grey.colors))
pairs(hec, highlighting = 2, diag_panel = pairs_diagonal_mosaic,
          diag_panel_args = list(fill = grey.colors, alternate_labels =TRUE))
```

plot.loglm	*Visualize Fitted Log-linear Models*

Description

Visualize fitted "loglm" objects by mosaic or association plots.

Usage

```
## S3 method for class 'loglm'
plot(x, panel = mosaic, type = c("observed", "expected"),
  residuals_type = c("pearson", "deviance"), gp = shading_hcl, gp_args = list(),
  ...)
```

Arguments

x	a fitted "loglm" object, see [loglm](#).
panel	a panel function for visualizing the observed values, residuals and expected values. Currently, [mosaic](#) and [assoc](#) in **vcd**.
type	a character string indicating whether the observed or the expected values of the table should be visualized.
residuals_type	a character string indicating the type of residuals to be computed.
gp	object of class "gpar", shading function or a corresponding generating function (see details and [shadings](#)). Ignored if shade = FALSE.
gp_args	list of arguments for the shading-generating function, if specified.
...	Other arguments passed to the panel function.

Details

The plot method for "loglm" objects by default visualizes the model using a mosaic plot (can be changed to an association plot by setting panel = assoc) with a shading based on the residuals of this model. The legend also reports the corresponding p value of the associated goodness-of-fit test. The mosaic and assoc methods are simple convenience interfaces to this plot method, setting the panel argument accordingly.

Value

The "structable" visualized is returned invisibly.

Author(s)

Achim Zeileis <Achim.Zeileis@R-project.org>

See Also

[loglm](#), [assoc](#), [mosaic](#), [strucplot](#)

Examples

```
library(MASS)
## mosaic display for PreSex model
data("PreSex")
fm <- loglm(~ PremaritalSex * ExtramaritalSex * (Gender + MaritalStatus),
  data = aperm(PreSex, c(3, 2, 4, 1)))
fm
## visualize Pearson statistic
plot(fm, split_vertical = TRUE)
## visualize LR statistic
plot(fm, split_vertical = TRUE, residuals_type = "deviance")

## conditional independence in UCB admissions data
data("UCBAdmissions")
fm <- loglm(~ Dept * (Gender + Admit), data = aperm(UCBAdmissions))

## use mosaic display
plot(fm, labeling_args = list(abbreviate = c(Admit = 3)))

## and association plot
plot(fm, panel = assoc)
assoc(fm)
```

PreSex   *Pre-marital Sex and Divorce*

Description

Data from Thornes \& Collard (1979), reported in Gilbert (1981), on pre- and extra-marital sex and divorce.

Usage

```
data("PreSex")
```

Format

A 4-dimensional array resulting from cross-tabulating 1036 observations on 4 variables. The variables and their levels are as follows:

No	Name	Levels
1	MaritalStatus	Divorced, Married
2	ExtramaritalSex	Yes, No
3	PremaritalSex	Yes, No
4	Gender	Women, Men

Source

Michael Friendly (2000), Visualizing Categorical Data: http://euclid.psych.yorku.ca/ftp/sas/vcd/catdata/marital.sas

References

G. N. Gilbert (1981), *Modelling Society: An Introduction to Loglinear Analysis for Social Researchers*. Allen and Unwin, London.

B. Thornes \& J. Collard (1979), *Who Divorces?*. Routledge \& Kegan, London.

M. Friendly (2000), *Visualizing Categorical Data*. SAS Institute, Cary, NC.

Examples

```
data("PreSex")

## Mosaic display for Gender and Premarital Sexual Experience
## (Gender Pre)
mosaic(margin.table(PreSex, c(3,4)),
                main = "Gender and Premarital Sex")

## (Gender Pre)(Extra)
mosaic(margin.table(PreSex, c(2,3,4)),
       expected = ~Gender * PremaritalSex + ExtramaritalSex ,
   main = "PreMaritalSex*Gender +Sex")

## (Gender Pre Extra)(Marital)
mosaic(PreSex,
       expected = ~Gender*PremaritalSex*ExtramaritalSex + MaritalStatus,
       main = "PreMarital*ExtraMarital + MaritalStatus")

## (GPE)(PEM)
mosaic(PreSex,
       expected = ~ Gender * PremaritalSex * ExtramaritalSex
                    + MaritalStatus * PremaritalSex * ExtramaritalSex,
       main = "G*P*E + P*E*M")
```

Punishment	*Corporal Punishment Data*

Description

Data from a study of the Gallup Institute in Denmark in 1979 about the attitude of a random sample of 1,456 persons towards corporal punishment of children.

Usage

```
data("Punishment")
```

Format

A data frame with 36 observations and 5 variables.

Freq frequency.

attitude factor indicating attitude: (no, moderate) punishment of children.

memory factor indicating whether the person had memories of corporal punishment as a child (yes, no).

education factor indicating highest level of education (elementary, secondary, high).

age factor indicating age group in years (15-24, 25-39, 40-).

Note

Anderson (1991) erroneously indicates the total sum of respondents to be 783.

Source

E. B. Andersen (1991), The Statistical Analysis of Categorical Data, pages 207–208.

References

E. B. Andersen (1991), *The Statistical Analysis of Categorical Data*. 2nd edition. Springer-Verlag, Berlin.

Examples

```
data("Punishment", package = "vcd")
pun <- xtabs(Freq ~ memory + attitude + age + education, data = Punishment)

## model: ~ (memory + attitude) * age * education
## use maximum sum-of-squares test/shading
cotabplot(~ memory + attitude | age + education, data = pun, panel = cotab_coindep,
  n = 5000, type = "assoc", test = "maxchisq", interpolate = 1:2)
```

RepVict   *Repeat Victimization Data*

Description

Data from Reiss (1980) given by Fienberg (1980) about instances of repeat victimization for households in the U.S. National Crime Survey.

Usage

```
data("RepVict")
```

Format

A 2-dimensional array resulting from cross-tabulating victimization. The variables and their levels are as follows:

No	Name	Levels
1	First Victimization	Rape, Assault, Robbery, Pickpocket, Personal Larceny, Burglary, Household Larceny, Auto Theft
2	Second Victimization	Rape, Assault, Robbery, Pickpocket, Personal Larceny, Burglary, Household Larceny, Auto Theft

Source

Michael Friendly (2000), Visualizing Categorical Data, page 113.

References

S. E. Fienberg (1980), *The Analysis of Cross-Classified Categorical Data*, MIT Press, Cambridge, 2nd edition.

A. J. J. Reiss (1980), Victim proneness by type of crime in repeat victimization. In S. E. Fienberg & A. J. J. Reiss (eds.), *Indicators of Crime and Criminal Justice*. U.S. Government Printing Office, Washington, DC.

M. Friendly (2000), *Visualizing Categorical Data*. SAS Institute, Cary, NC.

Examples

```
data("RepVict")

mosaic(RepVict[-c(4,7),-c(4,7)], gp = shading_max,
       main = "Repeat Victimization Data")
```

Rochdale	*Rochdale Data*

Description

Information on 665 households of Rochdale, Lancashire, UK. The study was conducted to identify influence factors on economical activity of wives.

Usage

```
data("Rochdale")
```

Format

A 8-dimensional array resulting from cross-tabulating 665 observations on 8 variables. The variables and their levels are as follows:

No	Name	Levels
1	EconActive	yes, no
2	Age	<38, >38
3	HusbandEmployed	yes, no
4	Child	yes, no

5	Education	yes, no
6	HusbandEducation	yes, no
7	Asian	yes, no
8	HouseholdWorking	yes, no

Note

Many observations are missing: only 91 out of all 256 combinations contain information.

Source

Whittaker (1990).

References

H. Hofmann (2003). Constructing and reading mosaicplots. *Computational Statistics & Data Analysis*, **43**, 4, 565–580.

J. Whittaker (1990), *Graphical Models on Applied Multivariate Statistics*, Wiley, New York.

Examples

```
data("Rochdale")
mosaic(Rochdale)
```

---

rootogram   *Rootograms*

Description

Rootograms of observed and fitted values.

Usage

```
## Default S3 method:
rootogram(x, fitted, names = NULL, scale = c("sqrt", "raw"),
  type = c("hanging", "standing", "deviation"),
  rect_gp = gpar(fill = "lightgray"), lines_gp = gpar(col = "red"),
  points_gp = gpar(col = "red"), pch = 19,
  xlab = NULL, ylab = NULL, ylim = NULL,
  name = "rootogram", newpage = TRUE, pop = TRUE, ...)
```

Arguments

x	either a vector or a 1-way table of frequencies.
fitted	a vector of fitted frequencies.
names	a vector of names passed to [grid_barplot](#), if set to NULL the names of x are used.
scale	a character string indicating whether the values should be plotted on the raw or square root scale.
type	a character string indicating if the bars for the observed frequencies should be hanging or standing or indicate the deviation between observed and fitted frequencies.
rect_gp	a "gpar" object controlling the grid graphical parameters of the rectangles.
lines_gp	a "gpar" object controlling the grid graphical parameters of the lines.
points_gp	a "gpar" object controlling the grid graphical parameters of the points.
pch	plotting character for the points.
xlab	a label for the x axis.
ylab	a label for the y axis.
ylim	limits for the y axis.
name	name of the plotting viewport.
newpage	logical. Should [grid.newpage](#) be called before plotting?
pop	logical. Should the viewport created be popped?
...	further arguments passed to [grid_barplot](#).

Details

The observed frequencies are displayed as bars and the fitted frequencies as a line. By default a sqrt scale is used to make the smaller frequencies more visible.

Author(s)

Achim Zeileis <Achim.Zeileis@R-project.org>

References

J. W. Tukey (1977), *Exploratory Data Analysis*. Addison Wesley, Reading, MA.

M. Friendly (2000), *Visualizing Categorical Data*. SAS Institute, Cary, NC.

See Also

[grid_barplot](#)

Examples

```
## Simulated data examples:
dummy <- rnbinom(200, size = 1.5, prob = 0.8)
observed <- table(dummy)
fitted1 <- dnbinom(as.numeric(names(observed)),
                   size = 1.5, prob = 0.8) * sum(observed)
fitted2 <- dnbinom(as.numeric(names(observed)),
                   size = 2, prob = 0.6) * sum(observed)
rootogram(observed, fitted1)
rootogram(observed, fitted2)

## Real data examples:
data("HorseKicks")
HK.fit <- goodfit(HorseKicks)
summary(HK.fit)
plot(HK.fit)
## or equivalently
rootogram(HK.fit)

data("Federalist")
F.fit <- goodfit(Federalist, type = "nbinomial")
summary(F.fit)
plot(F.fit)
```

Saxony   *Families in Saxony*

Description

Data from Geissler, cited in Sokal & Rohlf (1969) and Lindsey (1995) on gender distributions in families in Saxony in the 19th century.

Usage

```
data("Saxony")
```

Format

A 1-way table giving the number of male children in 6115 families of size 12. The variable and its levels are

No	Name	Levels
1	nMales	0, 1, ..., 12

Source

M. Friendly (2000), Visualizing Categorical Data, pages 40–42.

References

J. K. Lindsey (1995), *Analysis of Frequency and Count Data*. Oxford University Press, Oxford, UK.

R. R. Sokal & F. J. Rohlf (1969), *Biometry. The Principles and Practice of Statistics*. W. H. Freeman, San Francisco, CA.

M. Friendly (2000), *Visualizing Categorical Data*. SAS Institute, Cary, NC.

Examples

```
data("Saxony")
gf <- goodfit(Saxony, type = "binomial")
summary(gf)
plot(gf)
```

SexualFun   *Sex is Fun*

Description

Data from Hout et al. (1987) given by Agresti (1990) summarizing the responses of married couples to the questionnaire item: Sex is fun for me and my partner: (a) never or occasionally, (b) fairly often, (c) very often, (d) almost always.

Usage

```
data("SexualFun")
```

Format

A 2-dimensional array resulting from cross-tabulating the ratings of 91 married couples. The variables and their levels are as follows:

No	Name	Levels
1	Husband	Never Fun, Fairly Often, Very Often, Always Fun
2	Wife	Never Fun, Fairly Often, Very Often, Always Fun

Source

M. Friendly (2000), Visualizing Categorical Data, page 91.

References

A. Agresti (1990), *Categorical Data Analysis*. Wiley-Interscience, New York.

M. Hout, O. D. Duncan, M. E. Sobel (1987), Association and heterogeneity: Structural models of similarities and differences, *Sociological Methodology*, **17**, 145-184.

M. Friendly (2000), *Visualizing Categorical Data*. SAS Institute, Cary, NC.

Examples

```
data("SexualFun")

## Kappa statistics
Kappa(SexualFun)

## Agreement Chart
agreementplot(t(SexualFun), weights = 1)
## Partial Agreement Chart and B-Statistics
agreementplot(t(SexualFun),
              xlab = "Husband's Rating",
              ylab = "Wife's Rating",
              main = "Husband's and Wife's Sexual Fun")
```

---

shadings	*Shading-generating Functions for Residual-based Shadings*

Description

Shading-generating functions for computing residual-based shadings for mosaic and association plots.

Usage

```
shading_hcl(observed, residuals = NULL, expected = NULL, df = NULL,
  h = NULL, c = NULL, l = NULL, interpolate = c(2, 4), lty = 1,
  eps = NULL, line_col = "black", p.value = NULL, level = 0.95, ...)

shading_hsv(observed, residuals = NULL, expected = NULL, df = NULL,
  h = c(2/3, 0), s = c(1, 0), v = c(1, 0.5),
  interpolate = c(2, 4), lty = 1, eps = NULL, line_col = "black",
  p.value = NULL, level = 0.95, ...)

shading_max(observed = NULL, residuals = NULL, expected = NULL, df = NULL,
  h = NULL, c = NULL, l = NULL, lty = 1, eps = NULL, line_col = "black",
  level = c(0.9, 0.99), n = 1000, ...)

shading_Friendly(observed = NULL, residuals = NULL, expected = NULL, df = NULL,
  h = c(2/3, 0), lty = 1:2, interpolate = c(2, 4),
  eps = 0.01, line_col = "black", ...)

shading_sieve(observed = NULL, residuals = NULL, expected = NULL, df = NULL,
  h = c(260, 0), lty = 1:2, interpolate = c(2, 4),
  eps = 0.01, line_col = "black", ...)

shading_binary(observed = NULL, residuals = NULL, expected = NULL, df = NULL,
  col = NULL)
```

```
hcl2hex(h = 0, c = 35, l = 85, fixup = TRUE)
```

Arguments

observed	contingency table of observed values
residuals	contingency table of residuals
expected	contingency table of expected values
df	degrees of freedom of the associated independence model.
h	hue value in the HCL or HSV color description, has to be in [0, 360] for HCL and in [0, 1] for HSV colors. The default is to use blue and red for positive and negative residuals respectively. In the HCL specification it is `c(260, 0)` by default and for HSV `c(2/3, 0)`.
c	chroma value in the HCL color description. This controls the maximum chroma for significant and non-significant results respectively and defaults to `c(100, 20)`.
l	luminance value in the HCL color description. Defaults to `c(90, 50)` for small and large residuals respectively.
s	saturation value in the HSV color description. Defaults to `c(1, 0)` for large and small residuals respectively.
v	saturation value in the HSV color description. Defaults to `c(1, 0.5)` for significant and non-significant results respectively.
interpolate	a specification for mapping the absolute size of the residuals to a value in [0, 1]. This can be either a function or a numeric vector. In the latter case, a step function with steps of equal size going from 0 to 1 is used.
lty	a vector of two line types for positive and negative residuals respectively. Recycled if necessary.
eps	numeric tolerance value below which absolute residuals are considered to be zero, which is used for coding the border color and line type. If set to NULL (default), all borders have the default color specified by `line_col`. If set to a numeric value, all border colors corresponding to residuals with a larger absolute value are set to the full positive or negative color, respectively; borders corresponding to smaller residuals are are drawn with `line_col` and `lty[1]`. This is used principally in `shading_Friendly`.
line_col	default border color (for `shading_sieve`: default sieve color).
p.value	the p value associated with the independence model. By default, this is computed from a Chi-squared distribution with df degrees of freedom. `p.value` can be either a scalar or a `function(observed, residuals, expected, df)` that computes the p value from the data. If set to NA no inference is performed.
level	confidence level of the test used. If `p.value` is smaller than 1 - `level`, bright colors are used, otherwise dark colors are employed. For `shading_max` a vector of levels can be supplied. The corresponding critical values are then used as `interpolate` cut-offs.
n	number of permutations used in the call to `coindep_test`.
col	a vector of two colors for positive and negative residuals respectively.

| fixup | logical. Should the color be corrected to a valid RGB value before correction? |
| ... | Other arguments passed to hcl2hex or hsv, respectively. |

Details

These shading-generating functions can be passed to strucplot to generate residual-based shadings for contingency tables. strucplot calls these functions with the arguments observed, residuals, expected, df which give the observed values, residuals, expected values and associated degrees of freedom for a particular contingency table and associated independence model.

The shadings shading_hcl and shading_hsv do the same thing conceptually, but use HCL or HSV colors respectively. The former is usually preferred because they are perceptually based. Both shadings visualize the *sign* of the residuals of an independence model using two hues (by default: blue and red). The *absolute size* of the residuals is visualized by the colorfulness and the amount of grey, by default in three categories: very colorful for large residuals (> 4), less colorful for medium sized residuals (< 4 and > 2), grey/white for small residuals (< 2). More categories or a continuous scale can be specified by setting interpolate. Furthermore, the result of a significance test can be visualized by the amount of grey in the colors. If significant, a colorful palette is used, if not, the amount of color is reduced. See Zeileis, Meyer, and Hornik (2007) and diverge_hcl for more details.

The shading shading_max is applicable in 2-way contingency tables and uses a similar strategy as shading_hcl. But instead of using the cut-offs 2 and 4, it employs the critical values for the maximum statistic (by default at 90% and 99%). Consequently, color in the plot signals a significant result at 90% or 99% significance level, respectively. The test is carried out by calling coindep_test.

The shading shading_Friendly is very similar to shading_hsv, but additionally codes the sign of the residuals by different line types. See Friendly (1994) for more details. shading_sieve is similar, but uses HCL colors.

The shading shading_binary just visualizes the sign of the residuals by using two different colors (default: blue HCL(260, 50, 70) and red HCL(0, 50, 70)).

The color implementations employed are hsv from base R and polarLUV from the **colorspace** package, respectively. To transform the HCL coordinates to a hexadecimal color string (as returned by hsv), the function hex is employed. A convenience wrapper hcl2hex is provided.

Value

A shading function which takes only a single argument, interpreted as a vector/table of residuals, and returns a "gpar" object with the corresponding vector(s)/table(s) of graphical parameter(s).

Author(s)

Achim Zeileis <Achim.Zeileis@R-project.org>

References

Friendly M. (1994), Mosaic Displays for Multi-Way Contingency Tables. *Journal of the American Statistical Association*, **89**, 190–200.

Meyer D., Zeileis A., and Hornik K. (2006), The Strucplot Framework: Visualizing Multi-Way Contingency Tables with **vcd**. *Journal of Statistical Software*, **17**(3), 1–48. URL http://www.jstatsoft.org/v17/i03/. See also vignette("strucplot", package = "vcd").

Zeileis A., Meyer D., Hornik K. (2007), Residual-Based Shadings for Visualizing (Conditional) Independence. *Journal of Computational and Graphical Statistics*, **16**, 507–525.

Zeileis A., Hornik K. and Murrell P. (2008), Escaping RGBland: Selecting Colors for Statistical Graphics. *Computational Statistics & Data Analysis*, Forthcoming. Preprint available from `http://statmath.wu-wien.ac.at/~zeileis/papers/Zeileis+Hornik+Murrell-2008.pdf`.

See Also

`hex`, `polarLUV`, `hsv`, `mosaic`, `assoc`, `strucplot`, `diverge_hcl`

Examples

```
## load Arthritis data
data("Arthritis")
art <- xtabs(~Treatment + Improved, data = Arthritis)

## plain mosaic display without shading
mosaic(art)

## with shading for independence model
mosaic(art, shade = TRUE)
## which uses the HCL shading
mosaic(art, gp = shading_hcl)
## the residuals are two small to have color,
## hence the cut-offs can be modified
mosaic(art, gp = shading_hcl, gp_args = list(interpolate = c(1, 1.8)))
## the same with the Friendly palette
## (without significance testing)
mosaic(art, gp = shading_Friendly, gp_args = list(interpolate = c(1, 1.8)))

## assess independence using the maximum statistic
## cut-offs are now critical values for the test statistic
mosaic(art, gp = shading_max)

## association plot with shading as in base R
assoc(art, gp = shading_binary(col = c(1, 2)))
```

---

sieve   *Extended Sieve Plots*

Description

(Extended) sieve displays for n-way contingency tables: plots rectangles with areas proportional to the expected cell frequencies and filled with a number of squares equal to the observed frequencies. Thus, the densities visualize the deviations of the observed from the expected values.

Usage

```
## Default S3 method:
sieve(x, condvars = NULL, gp = NULL, shade = NULL,
  legend = FALSE, split_vertical = NULL, direction = NULL, spacing = NULL,
  spacing_args = list(), sievetype = c("observed","expected"),
  gp_tile = gpar(), main = NULL, sub = NULL, ...)
## S3 method for class 'formula'
sieve(formula, data, ..., main = NULL, sub = NULL, subset = NULL)
```

Arguments

x	a contingency table in array form, with optional category labels specified in the `dimnames(x)` attribute.
condvars	vector of integers or character strings indicating conditioning variables, if any. The table will be permuted to order them first.
formula	a formula specifying the variables used to create a contingency table from `data`. For convenience, conditioning formulas can be specified; the conditioning variables will then be used first for splitting. Formulas for sieve displays (unlike those for doubledecker plots) have no response variable.
data	either a data frame, or an object of class "table" or "ftable".
subset	an optional vector specifying a subset of observations to be used.
shade	logical specifying whether gp should be used or not (see gp). If TRUE and `expected` is unspecified, a default model is fitted: if condvars is specified, a corresponding conditional independence model, and else the total independence model. If shade is NULL (default), gp is used if specified.
sievetype	logical indicating whether rectangles should be filled according to `observed` or `expected` frequencies.
gp	object of class "gpar", shading function or a corresponding generating function (see details of [strucplot](#) and [shadings](#)). Components of "gpar" objects are recycled as needed along the last splitting dimension. The default is a modified version of [shading_Friendly](#): if sievetype is "observed", cells with positive residuals are painted with a red sieve, and cells with negative residuals with a blue one. If sievetype is "expected", the sieves' color is gray. Ignored if shade = FALSE.
gp_tile	object of class "gpar", controlling the appearance of all *static* elements of the cells (e.g., border and fill color).
legend	either a legend-generating function, a legend function (see details of [strucplot](#) and [legends](#)), or a logical value. If legend is NULL or TRUE and gp is a function, legend defaults to [legend_resbased](#).
split_vertical	vector of logicals of length k, where k is the number of margins of x (default: FALSE). Values are recycled as needed. A TRUE component indicates that the tile(s) of the corresponding dimension should be split vertically, FALSE means horizontal splits. Ignored if direction is not NULL.
direction	character vector of length k, where k is the number of margins of x (values are recycled as needed). For each component, a value of "h" indicates that the

	tile(s) of the corresponding dimension should be split horizontally, whereas "v" indicates vertical split(s).
spacing	spacing object, spacing function, or corresponding generating function (see strucplot for more information). The default is no spacing at all if x has two dimensions, and spacing_increase for more dimensions.
spacing_args	list of arguments for the generating function, if specified (see strucplot for more information).
main, sub	either a logical, or a character string used for plotting the main (sub) title. If logical and TRUE, the name of the data object is used.
...	Other arguments passed to strucplot

Details

sieve is a generic function which currently has a default method and a formula interface. Both are high-level interfaces to the strucplot function, and produce (extended) sieve displays. Most of the functionality is described there, such as specification of the independence model, labeling, legend, spacing, shading, and other graphical parameters.

The layout is very flexible: the specification of shading, labeling, spacing, and legend is modularized (see strucplot for details).

Value

The "structable" visualized is returned invisibly.

Note

To be faithful to the original definition by Riedwyl & Schüpbach, the default is to have no spacing between the tiles for two-way tables.

Author(s)

David Meyer <David.Meyer@R-project.org>

References

H. Riedwyl & M. Schüpbach (1994), Parquet diagram to plot contingency tables. In F. Faulbaum (ed.), *Softstat '93: Advances in Statistical Software*, 293–299. Gustav Fischer, New York.

M. Friendly (2000), Visualizing Categorical Data, SAS Institute, Cary, NC.

David Meyer, Achim Zeileis, and Kurt Hornik (2006). The strucplot framework: Visualizing multi-way contingency tables with **vcd**. *Journal of Statistical Software*, **17(3)**, 1-48. URL http://www.jstatsoft.org/v17/i03/ and available as vignette("strucplot").

See Also

assoc, strucplot, mosaic, structable, doubledecker

Examples

```
data("HairEyeColor")

## aggregate over 'sex':
(haireye <- margin.table(HairEyeColor, c(2,1)))

## plot expected values:
sieve(haireye, sievetype = "expected", shade = TRUE)

## plot observed table:
sieve(haireye, shade = TRUE)

## plot complete diagram:
sieve(HairEyeColor, shade = TRUE)

## example with observed values in the cells:
sieve(haireye, shade = TRUE, labeling = labeling_values,
               gp_text = gpar(fontface = 2))

## example with expected values in the cells:
sieve(haireye, shade = TRUE, labeling = labeling_values,
               value_type = "expected", gp_text = gpar(fontface = 2))

## an example for the formula interface:
data("VisualAcuity")
sieve(Freq ~ right + left,  data = VisualAcuity)
```

SpaceShuttle   *Space Shuttle O-ring Failures*

Description

Data from Dalal et al. (1989) about O-ring failures in the NASA space shuttle program. The damage index comes from a discussion of the data by Tufte (1997).

Usage

```
data("SpaceShuttle")
```

Format

A data frame with 24 observations and 6 variables.

FlightNumber Number of space shuttle flight.
Temperature temperature during start (in degrees F).
Pressure pressure.
Fail did any O-ring failures occur? (no, yes).
nFailures how many (of six) 0-rings failed?.
Damage damage index.

Source

Michael Friendly (2000), Visualizing Categorical Data: http://euclid.psych.yorku.ca/ftp/sas/vcd/catdata/orings.sas

References

S. Dalal, E. B. Fowlkes, B. Hoadly (1989), Risk analysis of the space shuttle: Pre-Challenger prediction of failure, *Journal of the American Statistical Association*, **84**, 945–957.

E. R. Tufte (1997), *Visual Explanations: Images and Quantities, Evidence and Narrative*. Graphics Press, Cheshire, CT.

M. Friendly (2000), *Visualizing Categorical Data*. SAS Institute, Cary, NC.

Examples

```
data("SpaceShuttle")
plot(nFailures/6 ~ Temperature, data = SpaceShuttle,
     xlim = c(30, 81), ylim = c(0,1),
     main = "NASA Space Shuttle O-Ring Failures",
     ylab = "Estimated failure probability",
     pch = 19, col = 4)
fm <- glm(cbind(nFailures, 6 - nFailures) ~ Temperature,
          data = SpaceShuttle,
          family = binomial)
lines(30 : 81,
      predict(fm, data.frame(Temperature = 30 : 81), type = "re"),
      lwd = 2)
abline(v = 31, lty = 3)
```

---

spacings	*Spacing-generating Functions*

Description

These functions generate spacing functions to be used with strucplot to obtain customized spaces between the elements of a strucplot.

Usage

```
spacing_equal(sp = unit(0.3, "lines"))
spacing_dimequal(sp)
spacing_increase(start = unit(0.3, "lines"), rate = 1.5)
spacing_conditional(sp = unit(0.3, "lines"), start = unit(2, "lines"), rate = 1.8)
spacing_highlighting(start = unit(0.2, "lines"), rate = 1.5)
```

Arguments

start   object of class "unit" indicating the start value for increasing spacings.

rate   increase rate for spacings.

sp   object of class "unit" specifying a fixed spacing.

Details

These generating functions return a function used by [strucplot](#) to generate appropriate spaces between tiles of a strucplot, using the dimnames information of the visualized table.

spacing_equal allows to specify one fixed space for *all* dimensions.

spacing_dimequal allows to specify a fixed space for *each* dimension.

spacing_increase creates increasing spaces for all dimensions, based on a starting value and an increase rate.

spacing_conditional combines spacing_equal and spacing_increase to create fixed spaces for conditioned dimensions, and increasing spaces for conditioning dimensions.

spacing_highlighting is essentially spacing_conditional but with the space of the last dimension set to 0. With a corresponding color scheme, this gives the impression of the last class being 'highlighted' in the penultimate class (as, e.g., in [doubledecker](#) plots).

Value

A spacing function with arguments:

d   "dim" attribute of a contingency table.

condvars   index vector of conditioning dimensions (currently only used by spacing_conditional).

This function computes a list of objects of class "unit". Each list element contains the spacing information for the corresponding dimension of the table. The length of the "unit" objects is $k - 1$, k number of levels of the corresponding factor.

Author(s)

David Meyer <David.Meyer@R-project.org>

References

Meyer, D., Zeileis, A., and Hornik, K. (2006), The strucplot framework: Visualizing multi-way contingency tables with **vcd**. *Journal of Statistical Software*, **17(3)**, 1-48. URL http://www.jstatsoft.org/v17/i03/ and available as vignette("strucplot").

See Also

[strucplot](#), [doubledecker](#)

Examples

```
data("Titanic")
strucplot(Titanic, spacing = spacing_increase(start = 0.5, rate = 1.5))
strucplot(Titanic, spacing = spacing_equal(1))
strucplot(Titanic, spacing = spacing_dimequal(1:4 / 4))
strucplot(Titanic, spacing = spacing_highlighting,
                   gp = gpar(fill = c("light gray","dark gray")))
data("PreSex")
strucplot(aperm(PreSex, c(1,4,2,3)), spacing = spacing_conditional,
          condvars = 2)
```

---

spine   *Spine Plots and Spinograms*

Description

Spine plots are a special cases of mosaic plots, and can be seen as a generalization of stacked (or highlighted) bar plots. Analogously, spinograms are an extension of histograms.

Usage

```
spine(x, ...)
## Default S3 method:
spine(x, y = NULL,
  breaks = NULL, ylab_tol = 0.05, off = NULL,
 main = "", xlab = NULL, ylab = NULL, ylim = c(0, 1), margins = c(5.1, 4.1, 4.1, 3.1),
  gp = gpar(), name = "spineplot", newpage = TRUE, pop = TRUE,
  ...)
## S3 method for class 'formula'
spine(formula, data = list(),
  breaks = NULL, ylab_tol = 0.05, off = NULL,
 main = "", xlab = NULL, ylab = NULL, ylim = c(0, 1), margins = c(5.1, 4.1, 4.1, 3.1),
  gp = gpar(), name = "spineplot", newpage = TRUE, pop = TRUE,
  ...)
```

Arguments

x	an object, the default method expects either a single variable (interpreted to be the explanatory variable) or a 2-way table. See details.
y	a "factor" interpreted to be the dependent variable
formula	a "formula" of type y ~ x with a single dependent "factor" and a single explanatory variable.
data	an optional data frame.
breaks	if the explanatory variable is numeric, this controls how it is discretized. breaks is passed to [hist](#) and can be a list of arguments.

ylab_tol	convenience tolerance parameter for y-axis annotation. If the distance between two labels drops under this threshold, they are plotted equidistantly.
off	vertical offset between the bars (in per cent). It is fixed to 0 for spinograms and defaults to 2 for spine plots.
main, xlab, ylab	
	character strings for annotation
ylim	limits for the y axis
margins	margins when calling `plotViewport`
gp	a `"gpar"` object controlling the grid graphical parameters of the rectangles. It should specify in particular a vector of `fill` colors of the same length as `levels(y)`. The default is to call `gray.colors`.
name	name of the plotting viewport.
newpage	logical. Should `grid.newpage` be called before plotting?
pop	logical. Should the viewport created be popped?
...	additional arguments passed to `plotViewport`.

Details

spine creates either a spinogram or a spine plot. It can be called via `spine(x, y)` or `spine(y ~ x)` where y is interpreted to be the dependent variable (and has to be categorical) and x the explanatory variable. x can be either categorical (then a spine plot is created) or numerical (then a spinogram is plotted). Additionally, `spine` can also be called with only a single argument which then has to be a 2-way table, interpreted to correspond to `table(x, y)`.

Spine plots are a generalization of stacked bar plots where not the heights but the widths of the bars corresponds to the relative frequencies of x. The heights of the bars then correspond to the conditional relative frequencies of y in every x group. This is a special case of a mosaic plot with specific spacing and shading.

Analogously, spinograms extend stacked histograms. As for the histogram, x is first discretized (using `hist`) and then for the discretized data a spine plot is created.

Value

The table visualized is returned invisibly.

Author(s)

Achim Zeileis <Achim.Zeileis@R-project.org>

References

Hummel, J. (1996), Linked bar charts: Analysing categorical data graphically. *Computational Statistics*, **11**, 23–33.

Hofmann, H., Theus, M. (2005), *Interactive graphics for visualizing conditional distributions*, Unpublished Manuscript.

See Also

cd_plot, mosaic, hist

Examples

```
## Arthritis data (dependence on a categorical variable)
data("Arthritis")
(spine(Improved ~ Treatment, data = Arthritis))

## Arthritis data (dependence on a numerical variable)
(spine(Improved ~ Age, data = Arthritis, breaks = 5))
(spine(Improved ~ Age, data = Arthritis, breaks = quantile(Arthritis$Age)))
(spine(Improved ~ Age, data = Arthritis, breaks = "Scott"))

## Space shuttle data (dependence on a numerical variable)
data("SpaceShuttle")
(spine(Fail ~ Temperature, data = SpaceShuttle, breaks = 3))
```

---

strucplot   *Structured Displays of Contingency Tables*

Description

This modular function visualizes certain aspects of high-dimensional contingency tables in a hier-archical way.

Usage

```
strucplot(x, residuals = NULL, expected = NULL,
  condvars = NULL, shade = NULL, type = c("observed", "expected"),
  residuals_type = NULL, df = NULL, split_vertical = NULL,
  spacing = spacing_equal, spacing_args = list(),
  gp = NULL, gp_args = list(),
  labeling = labeling_border, labeling_args = list(),
  core = struc_mosaic, core_args = list(),
  legend = NULL, legend_args = list(),
  main = NULL, sub = NULL, margins = unit(3, "lines"),
  title_margins = NULL, legend_width = NULL,
  main_gp = gpar(fontsize = 20), sub_gp = gpar(fontsize = 15),
  newpage = TRUE, pop = TRUE, keep_aspect_ratio = NULL, prefix = "", ...)
```

Arguments

x	a contingency table in array form, with optional category labels specified in the dimnames attribute.
residuals	optionally, an array of residuals of the same dimension as x (see details).

expected	optionally, an array of expected values of the same dimension as x, or alternatively the corresponding independence model specification as used by [loglin](#) or [loglm](#) (see details).
df	degrees of freedom passed to the shading functions used for inference. Will be calculated (and overwritten if specified) if both expected and residuals are NULL, or if expected is given a formula.
condvars	number of conditioning variables, if any; those are expected to be ordered first in the table. This information is used for computing the expected values, and is also passed to the spacing functions (see [spacings](#)).
shade	logical specifying whether gp should be used or not (see gp). If TRUE and expected is unspecified, a default model is fitted: if condvars is specified, a corresponding conditional independence model, and else the total independence model.
residuals_type	a character string indicating the type of residuals to be computed when none are supplied. If residuals is NULL, residuals_type must be one of "pearson" (default; giving components of Pearson's chi-squared), "deviance" (giving components of the likelihood ratio chi-squared), or "FT" for the Freeman-Tukey residuals. The value of this argument can be abbreviated. If residuals are specified, the value of residuals_type is just passed "as is" to the legend function.
type	a character string indicating whether the observed or the expected values of the table should be visualized.
split_vertical	vector of logicals of length k, where k is the number of margins of x (values are recycled as needed). A TRUE component indicates that the tile(s) of the corresponding dimension should be split vertically, FALSE means horizontal splits. Default is FALSE.
spacing	spacing object, spacing function, or a corresponding generating function (see details and [spacings](#)).
spacing_args	list of arguments for the spacing-generating function, if specified.
gp	object of class "gpar", shading function or a corresponding generating function (see details and [shadings](#)). Components of "gpar" objects are recycled as needed along the last splitting dimension. Ignored if shade = FALSE.
gp_args	list of arguments for the shading-generating function, if specified.
labeling	either a logical, or a labeling function, or a corresponding generating function (see details and [labelings](#). If FALSE or NULL, no labeling is produced.
labeling_args	list of arguments for the labeling-generating function, if specified.
core	either a core function, or a corresponding generating function (see details). Currently, generating functions for mosaic plots ([struc_mosaic](#)), association plots ([struc_assoc](#)), and sieve plots ([struc_sieve](#)) are provided.
core_args	list of arguments for the core-generating function, if specified.
legend	either a legend-generating function, or a legend function (see details and [legends](#)), or a logical. If legend is NULL or TRUE and gp is a function, legend defaults to [legend_resbased](#).
legend_args	list of arguments for the legend-generating function, if specified.

main	either a logical, or a character string used for plotting the main title. If main is a logical and TRUE, the name of the object supplied as x is used.
sub	a character string used for plotting the subtitle. If sub is a logical and TRUE and main is unspecified, the name of the object supplied as x is used.
margins	either an object of class "unit" of length 4, or a numeric vector of length 4. The elements are recycled as needed. The four components specify the top, right, bottom, and left margin of the plot, respectively. When a numeric vector is supplied, the numbers are interpreted as "lines" units. In addition, the unit or numeric vector may have named arguments ('top', 'right', 'bottom', and 'left'), in which case the non-named arguments specify the default values (recycled as needed), overloaded by the named arguments.
title_margins	either an object of class "unit" of length 2, or a numeric vector of length 2. The elements are recycled as needed. The two components specify the top and bottom *title* margin of the plot, respectively. The default for each *specified* title are 2 lines (and 0 else), except when a legend is plotted and keep_aspect_ratio is TRUE: in this case, the default values of both margins are set as to align the heights of legend and actual plot. When a numeric vector is supplied, the numbers are interpreted as "lines" units. In addition, the unit or numeric vector may have named arguments ('top' and 'bottom'), in which case the non-named argument specify the default value (recycled as needed), overloaded by the named arguments.
legend_width	An object of class "unit" of length 1 specifying the width of the legend (if any). Default: 5 lines.
pop	logical indicating whether the generated viewport tree should be removed at the end of the drawing or not.
main_gp, sub_gp	
	object of class "gpar" containing the graphical parameters used for the main (sub) title, if specified.
newpage	logical indicating whether a new page should be created for the plot or not.
keep_aspect_ratio	
	logical indicating whether the aspect ratio should be fixed or not. If unspecified, the default is TRUE for two-dimensional tables and FALSE otherwise.
prefix	optional character string used as a prefix for the generated viewport and grob names.
...	For convenience, list of arguments passed to the labeling-generating function used.

Details

This function—usually called by higher-level functions such as [assoc](#) and [mosaic](#)—generates conditioning plots of contingency tables. First, it sets up a set of viewports for main- and subtitles, legend, and the actual plot region. Then, residuals are computed as needed from observed and expected frequencies, where the expected frequencies are optionally computed for a specified independence model. Finally, the specified functions for spacing, gp, main plot, legend, and labeling are called to produce the plot. The function invisibly returns the "structable" object visualized.

Most elements of the plot, such as the core function, the spacing between the tiles, the shading of the tiles, the labeling, and the legend, are modularized in graphical appearance control ("grapcon") functions and specified as parameters. For each element *foo* (= spacing, labeling, core, or legend), strucplot takes two arguments: *foo* and *foo_args*, which can be used to specify the parameters in the following alternative ways:

1. Passing a suitable function to *foo* which subsequently will be called from strucplot to compute shadings, labelings, etc.

2. Passing a corresponding *generating* function to *foo*, along with parameters passed to *foo_args*, that generates such a function. Generating functions must inherit from classes "grapcon_generator" and "}*foo*\code{".

3. Except for the shading functions (*shading_bar*), passing *foo(foo_args)* to the *foo* argument.

4. For shadings and spacings, passing the final parameter object itself; see the corresponding help pages for more details on the data structures.

If legends are drawn, a 'cinemascope'-like layout is used for the plot to preserve the 1:1 aspect ratio.

If type = "expected", the expected values are passed to the observed argument of the core function, and the observed values to the expected argument.

Although the gp argument is typically used for shading, it can be used for arbitrary modifications of the tiles' graphics parameters (e.g., for highlighting particular cells, etc.).

Value

Invisibly, an object of class "structable" corresponding to the plot.

Note

The created viewports, as well as the tiles and bullets, are named and thus can conveniently modified after a plot has been drawn (and pop = FALSE).

Author(s)

David Meyer <David.Meyer@R-project.org>

References

Meyer D., Zeileis A., and Hornik K. (2006), The strucplot framework: Visualizing multi-way contingency tables with **vcd**. *Journal of Statistical Software*, **17(3)**, 1-48. URL http://www.jstatsoft.org/v17/i03/ and available as vignette("strucplot").

See Also

assoc, mosaic, sieve, struc_assoc, struc_sieve, struc_mosaic, structable, doubledecker, labelings, shadings, legends, spacings

Examples

```
data("Titanic")

strucplot(Titanic)
strucplot(Titanic, core = struc_assoc)
strucplot(Titanic, spacing = spacing_increase,
                    spacing_args = list(start = 0.5, rate = 1.5))
strucplot(Titanic, spacing = spacing_increase(start = 0.5, rate = 1.5))

## modify a tile's color
strucplot(Titanic, pop = FALSE)
grid.edit("rect:Class=1st,Sex=Male,Age=Adult,Survived=Yes",
          gp = gpar(fill = "red"))
```

---

structable   *Structured Contingency Tables*

Description

This function produces a 'flat' representation of a high-dimensional contingency table constructed by recursive splits (similar to the construction of mosaic displays).

Usage

```
## S3 method for class 'formula'
structable(formula, data,
direction = NULL, split_vertical = NULL, ..., subset, na.action)
## Default S3 method:
structable(..., direction = NULL, split_vertical = FALSE)
```

Arguments

formula	a formula object with possibly both left and right hand sides specifying the column and row variables of the flat table.
data	a data frame, list or environment containing the variables to be cross-tabulated, or an object inheriting from class table.
subset	an optional vector specifying a subset of observations to be used. Ignored if data is a contingency table.
na.action	a function which indicates what should happen when the data contain NAs. Ignored if data is a contingency table
...	R objects which can be interpreted as factors (including character strings), or a list (or data frame) whose components can be so interpreted, or a contingency table object of class "table" or "ftable".
split_vertical	logical vector indicating, for each dimension, whether it should be split vertically or not (default: FALSE). Values are recycled as needed. If the argument is of length 1, the value is alternated for all dimensions. Ignored if direction is provided.

direction	character vector alternatively specifying the splitting direction ("h" for horizontal and "v" for vertical splits). Values are recycled as needed. If the argument is of length 1, the value is alternated for all dimensions.

Details

This function produces textual representations of mosaic displays, and thus 'flat' contingency tables. The formula interface is quite similar to the one of `ftable`, but also accepts the `mosaic`-like formula interface (empty left-hand side). Note that even if the `ftable` interface is used, the `split_vertical` or `direction` argument is needed to specify the *order* of the horizontal and vertical splits. If pretabulated data with a `Freq` column is used, than the left-hand side should be left empty—the `Freq` column will be handled correctly.

"structable" objects can be subset using the `[` and `[[` operators, using either level indices or names (see examples). The corresponding replacement functions are available as well. In addition, appropriate `aperm`, `cbind`, `rbind`, `length`, `dim`, and `is.na` methods do exist.

Value

An object of class "structable", inheriting from class "ftable", with the splitting information ("split_vertical") as additional attribute.

Author(s)

David Meyer <David.Meyer@R-project.org>

References

Meyer, D., Zeileis, A., and Hornik, K. (2006), The strucplot framework: Visualizing multi-way contingency tables with **vcd**. *Journal of Statistical Software*, **17(3)**, 1-48. URL http://www.jstatsoft.org/v17/i03/ and available as vignette("strucplot").

See Also

`strucplot`, `mosaic`, `ftable`

Examples

```
structable(Titanic)
structable(Titanic, split_vertical = c(TRUE, TRUE, FALSE, FALSE))
structable(Titanic, direction = c("h","h","v","v"))
structable(Sex + Class ~ Survived + Age, data = Titanic)

## subsetting of structable objects
(hec <- structable(aperm(HairEyeColor)))

## The "[" operator treats structables as a block-matrix and selects parts of the matrix:
hec[1]
hec[2]
hec[1,c(2,4)]
hec["Male",c("Blue","Green")]
```

```
## replacement funcion:
tmp <- hec
(tmp[1,2:3] <- tmp[2,c(1,4)])

## In contrast, the "[[" operator treats structables as two-dimensional
## lists. Indexing conditions on specified levels and thus reduces the dimensionality:

## seek subtables conditioning on levels of the first dimension:
hec[[1]]
hec[[2]]

## Seek subtable from the first two dimensions, given the level "Male"
## of the first variable, and "Brown" from the second
## (the following two commands are equivalent):
hec[["Male"]][["Brown"]]
hec[[c("Male","Brown")]]

## Seeking subtables by conditioning on row and/or column variables:
hec[["Male","Hazel"]]
hec[[c("Male","Brown"),]]
hec[[c("Male","Brown"),"Hazel"]]

## a few other operations
t(hec)
dim(hec)
dimnames(hec)
as.matrix(hec)
length(hec)
cbind(hec[,1],hec[,3])

as.vector(hec) ## computed on the _multiway_ table
as.vector(unclass(hec))
```

---

struc_assoc	*Core-generating Function for Association Plots*

Description

Core-generating function for `strucplot` returning a function producing association plots.

Usage

```
struc_assoc(compress = TRUE, xlim = NULL, ylim = NULL,
  yspace = unit(0.5, "lines"), xscale = 0.9, gp_axis = gpar(lty = 3))
```

Arguments

compress   logical; if FALSE, the space between the rows (columns) are chosen such that
              the *total* heights (widths) of the rows (column) are all equal. If TRUE, the space
              between the rows and columns is fixed and hence the plot is more "compressed".

xlim   either a $2 \times k$ matrix of doubles, k the number of total columns of the plot, or
                  a recycled vector from which such a matrix will be constructed. The columns
                  of xlim correspond to the columns of the association plot, the rows describe
                  the column ranges (minimums in the first row, maximums in the second row).
                  If xlim is NULL, the ranges are determined from the residuals according to
                  compress (if TRUE: widest range from each column, if FALSE: from the whole
                  association plot matrix).

ylim   either a $2 \times k$ matrix of doubles, k the number of total rows of the plot, or a
                  recycled vector from which such a matrix will be constructed. The columns of
                  ylim correspond to the rows of the association plot, the rows describe the col-
                  umn ranges (minimums in the first row, maximums in the second row). If ylim
                  is NULL, the ranges are determined from the residuals according to compress (if
                  TRUE: widest range from each row, if FALSE: from the whole association plot
                  matrix).

xscale   scale factor resizing the tile's width, thus adding additional space between the
                  tiles.

yspace   object of class "unit" specifying additional space separating the rows.

gp_axis   object of class "gpar" specifying the visual aspects of the tiles' baseline.

Details

This function is usually called by strucplot (typically when called by assoc) and returns a func-
tion used by strucplot to produce association plots.

Value

A function with arguments:

residuals   table of residuals.

observed   not used by struc_assoc.

expected   table of expected frequencies.

spacing   object of class "unit" specifying the space between the tiles.

gp   list of gpar objects used for the drawing the tiles.

split_vertical   vector of logicals indicating, for each dimension of the table, the split direction.

Author(s)

David Meyer <David.Meyer@R-project.org>

References

Cohen, A. (1980), On the graphical display of the significant components in a two-way contingency
table. *Communications in Statistics—Theory and Methods*, **A9**, 1025–1041.

Friendly, M. (1992), Graphical methods for categorical data. *SAS User Group International Con-
ference Proceedings*, **17**, 190–200. http://datavis.ca/sugi/sugi17.pdf

Meyer, D., Zeileis, A., and Hornik, K. (2006), The strucplot framework: Visualizing multi-way con-
tingency tables with **vcd**. *Journal of Statistical Software*, **17(3)**, 1-48. URL http://www.jstatsoft.org/v17/i03/
and available as vignette("strucplot").

See Also

assoc, strucplot, structable

Examples

```
## UCB Admissions
data("UCBAdmissions")
ucb <- aperm(UCBAdmissions)

## association plot for conditional independence
strucplot(ucb, expected = ~ Dept * (Admit + Gender),
  core = struc_assoc(ylim = c(-4, 4)), labeling_args = list(abbreviate = c(Admit = 3)))
```

---

struc_mosaic	*Core-generating Function for Mosaic Plots*

Description

Core-generating function for strucplot returning a function producing mosaic plots.

Usage

```
struc_mosaic(zero_size = 0.5, zero_split = FALSE, zero_shade = TRUE,
             zero_gp = gpar(col = 0), panel = NULL)
```

Arguments

zero_size   size of the bullets used for zero-entries in the contingency table (if 0, no bullets are drawn).

zero_split   logical controlling whether zero cells should be further split. If FALSE and zero_shade is FALSE, only one bullet is drawn (centered) for unsplit zero cells. If FALSE and zero_shade is TRUE, a bullet for each zero cell is drawn to allow, e.g., residual-based shadings to be effective also for zero cells.

zero_shade   logical controlling whether zero bullets should be shaded.

zero_gp   object of class "gpar" used for zero bullets in case they are *not* shaded.

panel   Optional function with arguments: residuals, observed, expected, index, gp, and name called by the struc_mosaic workhorse for each tile that is drawn in the mosaic. index is an integer vector with the tile's coordinates in the contingency table, gp a gpar object for the tile, and name a label to be assigned to the drawn grid object.

Details

This function is usually called by strucplot (typically when called by mosaic) and returns a function used by strucplot to produce mosaic plots.

Value

A function with arguments:

residuals   table of residuals.

observed   table of observed values.

expected   not used by struc_mosaic.

spacing   object of class "unit" specifying the space between the tiles.

gp   list of gpar objects used for the drawing the tiles.

split_vertical   vector of logicals indicating, for each dimension of the table, the split direction.

Author(s)

David Meyer <David.Meyer@R-project.org>

References

Cohen, A. (1980), On the graphical display of the significant components in a two-way contingency table. *Communications in Statistics—Theory and Methods*, **A9**, 1025–1041.

Friendly, M. (1992), Graphical methods for categorical data. *SAS User Group International Conference Proceedings*, **17**, 190–200. http://datavis.ca/sugi/sugi17.pdf

Meyer, D., Zeileis, A., and Hornik, K. (2006), The strucplot framework: Visualizing multi-way contingency tables with **vcd**. *Journal of Statistical Software*, **17(3)**, 1-48. URL http://www.jstatsoft.org/v17/i03/ and available as vignette("strucplot").

See Also

mosaic, strucplot, structable

Examples

```
## Titanic data
data("Titanic")
## mosaic plot with large zeros
strucplot(Titanic, core = struc_mosaic(zero_size = 1))
```

---

struc_sieve   *Core-generating Function for Sieve Plots*

Description

Core-generating function for strucplot returning a function producing sieve plots.

Usage

```
struc_sieve(sievetype = c("observed","expected"), gp_tile = gpar())
```

Arguments

sievetype	logical indicating whether rectangles should be filled according to observed or expected frequencies.
gp_tile	object of class "gpar", controlling the appearance of all *static* elements of the cells (e.g., border and fill color).

Details

This function is usually called by [strucplot](#) (typically when called by [sieve](#)) and returns a function used by [strucplot](#) to produce sieve plots.

Value

A function with arguments:

residuals	table of residuals.
observed	table of observed values.
expected	not used by struc_sieve.
spacing	object of class "unit" specifying the space between the tiles.
gp	list of gpar objects used for the drawing the tiles.
split_vertical	vector of logicals indicating, for each dimension of the table, the split direction.

Author(s)

David Meyer <David.Meyer@R-project.org>

References

Riedwyl, H., and Schüpbach, M. (1994), Parquet diagram to plot contingency tables. In F. Faulbaum (ed.), *Softstat '93: Advances in Statistical Software*, 293–299. Gustav Fischer, New York.

Friendly, M. (2000), Visualizing Categorical Data, SAS Institute, Cary, NC.

Meyer, D., Zeileis, A., and Hornik, K. (2006), The strucplot framework: Visualizing multi-way contingency tables with **vcd**. *Journal of Statistical Software*, **17(3)**, 1-48. URL http://www.jstatsoft.org/v17/i03/ and available as vignette("strucplot").

See Also

[sieve](#), [strucplot](#), [structable](#)

Examples

```
## Titanic data
data("Titanic")
strucplot(Titanic, core = struc_sieve)
```

Suicide   *Suicide Rates in Germany*

Description

Data from Heuer (1979) on suicide rates in West Germany classified by age, sex, and method of suicide.

Usage

```
data("Suicide")
```

Format

A data frame with 306 observations and 6 variables.

Freq frequency of suicides.

sex factor indicating sex (male, female).

method factor indicating method used.

age age (rounded).

age.group factor. Age classified into 5 groups.

method2 factor indicating method used (same as method but some levels are merged).

Source

Michael Friendly (2000), Visualizing Categorical Data: http://euclid.psych.yorku.ca/ftp/sas/vcd/catdata/suicide.sas

References

J. Heuer (1979), *Selbstmord bei Kindern und Jugendlichen*. Ernst Klett Verlag, Stuttgart.

M. Friendly (2000), *Visualizing Categorical Data*. SAS Institute, Cary, NC.

Examples

```
data("Suicide")
structable(~ sex + method2 + age.group, data = Suicide)
```

---

table2d_summary	*Summary of a 2-way Table*

Description

Prints a 2-way contingency table along with percentages, marginal, and conditional distributions.

Usage

```
table2d_summary(object, margins = TRUE, percentages = FALSE,
  conditionals = c("none", "row", "column"), ...)
```

Arguments

object	a $r \times c$-contingency table
margins	if TRUE, marginal distributions are computed.
percentages	if TRUE, relative frequencies are computed.
conditionals	if not "none", the conditional distributions, given the row/column factor, are computed.
...	currently not used.

Value

Returns invisibly a $r \times c \times k$ table, k depending on the amount of choices (at most 3).

Author(s)

David Meyer <David.Meyer@R-project.org>

See Also

[mar_table](), [prop.table](), [independence_table]()

Examples

```
data("UCBAdmissions")
table2d_summary(margin.table(UCBAdmissions, 1:2))
```

---

ternaryplot   *Ternary Diagram*

Description

Visualizes compositional, 3-dimensional data in an equilateral triangle.

Usage

```
ternaryplot(x, scale = 1, dimnames = NULL,
  dimnames_position = c("corner","edge","none"),
  dimnames_color = "black", id = NULL, id_color = "black",
  id_just = c("center", "center"),
  coordinates = FALSE, grid = TRUE, grid_color = "gray",
  labels = c("inside", "outside", "none"),
  labels_color = "darkgray", border = "black", bg = "white",
  pch = 19, cex = 1, prop_size = FALSE, col = "red",
  main = "ternary plot", newpage = TRUE, pop = TRUE, ...)
```

Arguments

x	a matrix with three columns.
scale	row sums scale to be used.
dimnames	dimension labels (defaults to the column names of x).
dimnames_position, dimnames_color	
	position and color of dimension labels.
id	optional labels to be plotted below the plot symbols. coordinates and id are mutual exclusive.
id_color	color of these labels.
id_just	character vector of length 1 or 2 indicating the justification of these labels.
coordinates	if TRUE, the coordinates of the points are plotted below them. coordinates and id are mutual exclusive.
grid	if TRUE, a grid is plotted. May optionally be a string indicating the line type (default: "dotted").
grid_color	grid color.
labels, labels_color	
	position and color of the grid labels.
border	color of the triangle border.
bg	triangle background.
pch	plotting character. Defaults to filled dots.
cex	a numerical value giving the amount by which plotting text and symbols should be scaled relative to the default. Ignored for the symbol size if prop_size is not FALSE.

prop_size	if TRUE, the symbol size is plotted proportional to the row sum of the three variables, i.e., represents the weight of the observation.
col	plotting color.
main	main title.
newpage	if TRUE, the plot will appear on a new graphics page.
pop	logical; if TRUE, all newly generated viewports are popped after plotting.
...	additional graphics parameters (see par)

Details

A points' coordinates are found by computing the gravity center of mass points using the data entries as weights. Thus, the coordinates of a point $P(a, b, c)$, $a + b + c = 1$, are: $P(b + c/2, c\sqrt{3}/2)$.

Author(s)

David Meyer <David.Meyer@R-project.org>

References

M. Friendly (2000), *Visualizing Categorical Data*. SAS Institute, Cary, NC.

Examples

```
data("Arthritis")

## Build table by crossing Treatment and Sex
tab <- as.table(xtabs(~ I(Sex:Treatment) + Improved, data = Arthritis))

## Mark groups
col <- c("red", "red", "blue", "blue")
pch <- c(1, 19, 1, 19)

## plot
ternaryplot(
            tab,
            col = col,
            pch = pch,
            prop_size = TRUE,
            bg = "lightgray",
            grid_color = "white",
            labels_color = "white",
            main = "Arthritis Treatment Data"
          )
## legend
grid_legend(0.8, 0.7, pch, col, rownames(tab), title = "GROUP")

## Titanic
data("Lifeboats")
attach(Lifeboats)
ternaryplot(
```

```
    Lifeboats[,4:6],
    pch = ifelse(side == "Port", 1, 19),
    col = ifelse(side == "Port", "red", "blue"),
    id  = ifelse(men / total > 0.1, as.character(boat), NA),
    prop_size = 2,
    dimnames_position = "edge",
    main = "Lifeboats on Titanic"
)
grid_legend(0.8, 0.9, c(1, 19),
  c("red", "blue"), c("Port", "Starboard"),
  title = "SIDE")


## Hitters
data("Hitters")
attach(Hitters)
colors <- c("black","red","green","blue","red","black","blue")
pch <- substr(levels(Positions), 1, 1)
ternaryplot(
  Hitters[,2:4],
  pch = as.character(Positions),
  col = colors[as.numeric(Positions)],
  main = "Baseball Hitters Data"
)
grid_legend(0.8, 0.9, pch, colors, levels(Positions),
  title = "POSITION(S)")
```

---

tile   *Tile Plot*

Description

Plots a tile display.

Usage

```
## Default S3 method:
tile(x,
        tile_type = c("area", "squaredarea", "height", "width"),
        halign = c("left", "center", "right"),
        valign = c("bottom", "center", "top"),
        split_vertical = NULL,
        shade = FALSE,
        spacing = spacing_equal(unit(1, "lines")),
        set_labels = NULL,
        margins = unit(3, "lines"),
        keep_aspect_ratio = FALSE,
        legend_width = NULL,
```

```
          squared_tiles = TRUE,
        main = NULL, sub = NULL, ...)
## S3 method for class 'formula'
tile(formula, data,
  ..., main = NULL, sub = NULL, subset = NULL, na.action = NULL)
```

Arguments

x	a contingency table, or an object coercible to one.
formula	a formula specifying the variables used to create a contingency table from data.
data	either a data frame, or an object of class "table" or "ftable".
subset	an optional vector specifying a subset of observations to be used.
na.action	a function which indicates what should happen when the data contain NAs. Ignored if data is a contingency table.
tile_type	character string indicating how the tiles should reflect the table frequencies (see details).
halign, valign	character string specifying the horizontal and vertical alignment of the tiles.
split_vertical	vector of logicals of length k, where k is the number of margins of x (values are recycled as needed). A TRUE component indicates that the tile(s) of the corresponding dimension should be split vertically, FALSE means horizontal splits. Default is FALSE.
spacing	spacing object, spacing function, or corresponding generating function (see [strucplot](#) for more information).
set_labels	An optional character vector with named components replacing the so-specified variable names. The component names must exactly match the variable names to be replaced.
shade	logical specifying whether shading should be enabled or not (see [strucplot](#)).
margins	either an object of class "unit" of length 4, or a numeric vector of length 4. The elements are recycled as needed. The four components specify the top, right, bottom, and left margin of the plot, respectively. When a numeric vector is supplied, the numbers are interpreted as "lines" units. In addition, the unit or numeric vector may have named arguments ('top', 'right', 'bottom', and 'left'), in which case the non-named arguments specify the default values (recycled as needed), overloaded by the named arguments.
legend_width	An object of class "unit" of length 1 specifying the width of the legend (if any). Default: 5 lines.
keep_aspect_ratio	logical indicating whether the aspect ratio should be fixed or not. The default is FALSE to enable the creation of squared tiles.
squared_tiles	logical indicating whether white space should be added as needed to rows or columns to obtain squared tiles in case of an unequal number of row and column labels.
main, sub	either a logical, or a character string used for plotting the main (sub) title. If logical and TRUE, the name of the data object is used.
...	Other arguments passed to [strucplot](#)

Details

A tile plot is a matrix of tiles. For each tile, either the "width", "height", "area", or squared area is proportional to the corresponding entry. The first three options allow column-wise, row-wise and overall comparisons, respectively. The last variant allows to compare the tiles both column-wise and row-wise, considering either the width or the height, respectively.

In contrast to other high-level strucplot functions, tile also accepts a table with duplicated levels (see examples). In this case, artificial dimnames will be created, and the actual ones are drawn using set_labels.

Note that multiway-tables are first "flattened" using structable.

Value

The "structable" visualized is returned invisibly.

Author(s)

David Meyer <David.Meyer@R-project.org>

See Also

[assoc](#), [strucplot](#), [mosaic](#), [structable](#),

Examples

```
data("Titanic")

## default plot
tile(Titanic)
tile(Titanic, type = "expected")
tile(Titanic, shade = TRUE)

## some variations
tile(Titanic, tile_type = "squaredarea")
tile(Titanic, tile_type = "width", squared_tiles = FALSE)
tile(Titanic, tile_type = "height", squared_tiles = FALSE)
tile(Titanic, tile_type = "area", halign = "center", valign = "center")

## repeat levels
tile(Titanic[,,,c(1,2,1,2)])
```

Trucks   *Truck Accidents Data*

Description

Data from a study in England in two periods from November 1969 to October 1971 and November 1971 to October 1973. A new compulsory safety measure for trucks was introduced in October 1971. Therefore, the question is whether the safety measure had an effect on the number of accidents and on the point of collision on the truck.

Usage

```
data("Trucks")
```

Format

A data frame with 24 observations on 5 variables.

Freq frequency of accidents involving trucks.

period factor indicating time period (before, after) 1971-11-01.

collision factor indicating whether the collision was in the back or forward (including the front and the sides) of the truck (back, forward).

parked factor indicating whether the truck was parked (yes, no).

light factor indicating light conditions: day light (daylight), night on an illuminated road (night, illuminate), night on a dark road (night, dark).

Source

E. B. Andersen (1991), The Statistical Analysis of Categorical Data, Table 6.8.

References

E. B. Andersen (1991), *The Statistical Analysis of Categorical Data*. 2nd edition. Springer-Verlag, Berlin.

Examples

```
library(MASS)
data("Trucks")
tab <- xtabs(Freq ~ period + collision + light + parked, data = Trucks)
loglm(~ (collision + period) * parked * light, data = tab)
doubledecker(collision ~ parked + light + period, data = tab)
cotabplot(tab, panel = cotab_coindep)
```

UKSoccer   *UK Soccer Scores*

Description

Data from Lee (1997), on the goals scored by Home and Away teams in the Premier Football League, 1995/6 season.

Usage

```
data("UKSoccer")
```

Format

A 2-dimensional array resulting from cross-tabulating the number of goals scored in 380 games. The variables and their levels are as follows:

No	Name	Levels
1	Home	0, 1, ..., 4
2	Away	0, 1, ..., 4

Source

M. Friendly (2000), Visualizing Categorical Data, page 27.

References

A. J. Lee (1997), Modelling scores in the Premier League: Is Manchester United really the best?, *Chance*, **10**(1), 15–19.

M. Friendly (2000), *Visualizing Categorical Data*. SAS Institute, Cary, NC.

See Also

Bundesliga

Examples

```
data("UKSoccer")
mosaic(UKSoccer, gp = shading_max, main = "UK Soccer Scores")
```

VisualAcuity	*Visual Acuity in Left and Right Eyes*

Description

Data from Kendall & Stuart (1961) on unaided vision among 3,242 men and 7,477 women, all aged 30-39 and employed in the U.K. Royal Ordnance factories 1943-1946.

Usage

```
data("VisualAcuity")
```

Format

A data frame with 32 observations and 4 variables.

Freq frequency of visual acuity measurements.

right visual acuity on right eye.

left visual acuity on left eye.

gender factor indicating gender of patient.

Source

M. Friendly (2000), Visualizing Categorical Data: http://euclid.psych.yorku.ca/ftp/sas/vcd/catdata/vision.sas

References

M. G. Kendall & A. Stuart (1961), *The Advanced Theory of Statistics*, Vol. 2. Griffin, London.

M. Friendly (2000), *Visualizing Categorical Data*. SAS Institute, Cary, NC.

Examples

```
data("VisualAcuity")
structable(~ gender + left + right, data = VisualAcuity)
```

VonBort   *Von Bortkiewicz Horse Kicks Data*

Description

Data from von Bortkiewicz (1898), given by Andrews \& Herzberg (1985), on number of deaths by horse or mule kicks in 14 corps of the Prussian army.

Usage

```
data("VonBort")
```

Format

A data frame with 280 observations and 4 variables.

deaths number of deaths.

year year of the deaths.

corps factor indicating the corps.

fisher factor indicating whether the corresponding corps was considered by Fisher (1925) or not.

Source

Michael Friendly (2000), Visualizing Categorical Data: http://euclid.psych.yorku.ca/ftp/sas/vcd/catdata/vonbort.sas

References

D. F. Andrews \& A. M. Herzberg (1985), *Data: A Collection of Problems from Many Fields for the Student and Research Worker*. Springer-Verlag, New York, NY.

R. A. Fisher (1925), *Statistical Methods for Research Workers*. Oliver & Boyd, London.

L. von Bortkiewicz (1898), *Das Gesetz der kleinen Zahlen*. Teubner, Leipzig.

M. Friendly (2000), *Visualizing Categorical Data*. SAS Institute, Cary, NC.

See Also

HorseKicks for a popular subsample.

Examples

```
data("VonBort")
## HorseKicks data
xtabs(~ deaths, data = VonBort, subset = fisher == "yes")
```

WeldonDice   *Weldon's Dice Data*

Description

Data from Pearson (1900) about the frequency of 5s and 6s in throws of 12 dice. Weldon tossed the dice 26,306 times and reported his results in a letter to Francis Galton on 1894-02-02.

Usage

```
data("WeldonDice")
```

Format

A 1-way table giving the frequency of a 5 or a 6 in 26,306 throws of 12 dice where 10 indicates '10 or more' 5s or 6s. The variable and its levels are

No	Name	Levels
1	n56	0, 1, …, 10

Source

M. Friendly (2000), Visualizing Categorical Data, pages 20–21.

References

K. Pearson (1900), On the criterion that a given system of deviations from the probable in the case of a correlated system of variables is such that it can be reasonably supposed to have arisen by random sampling, *Philosophical Magazine*, **50** (5th series), 157–175.

M. Friendly (2000), *Visualizing Categorical Data*. SAS Institute, Cary, NC.

Examples

```
data("WeldonDice")
gf <- goodfit(WeldonDice, type = "binomial")
summary(gf)
plot(gf)
```

WomenQueue   *Women in Queues*

Description

Data from Jinkinson \& Slater (1981) and Hoaglin \& Tukey (1985) reporting the frequency distribution of females in 100 queues of length 10 in a London Underground station.

Usage

```
data("WomenQueue")
```

Format

A 1-way table giving the number of women in 100 queues of length 10. The variable and its levels are

No	Name	Levels
1	nWomen	0, 1, ..., 10

Source

M. Friendly (2000), Visualizing Categorical Data, pages 19–20.

References

D. C. Hoaglin \& J. W. Tukey (1985), Checking the shape of discrete distributions. In D. C. Hoaglin, F. Mosteller, J. W. Tukey (eds.), *Exploring Data Tables, Trends and Shapes*, chapter 9. John Wiley \& Sons, New York.

R. A. Jinkinson \& M. Slater (1981), Critical discussion of a graphical method for identifying discrete distributions, *The Statistician*, **30**, 239–248.

M. Friendly (2000), *Visualizing Categorical Data*. SAS Institute, Cary, NC.

Examples

```
data("WomenQueue")
gf <- goodfit(WomenQueue, type = "binomial")
summary(gf)
plot(gf)
```

---

woolf_test   *Woolf Test*

Description

Test for homogeneity on $2 \times 2 \times k$ tables over strata (i.e., whether the log odds ratios are the same in all strata).

Usage

```
woolf_test(x)
```

Arguments

x   A $2 \times 2 \times k$ table.

Value

A list of class `"htest"` containing the following components:

statistic	the chi-squared test statistic.
parameter	degrees of freedom of the approximate chi-squared distribution of the test statistic.
p.value	p-value for the test.
method	a character string indicating the type of test performed.
data.name	a character string giving the name(s) of the data.
observed	the observed counts.
expected	the expected counts under the null hypothesis.

References

Woolf, B. 1955. On estimating the relation between blood group and disease. *Ann. Human Genet.* (London) **19**, 251-253.

See Also

mantelhaen.test

Examples

```
data("CoalMiners")
woolf_test(CoalMiners)
```

Index